

Work with Big Data in R

Ideas and Advice



Garrett Grolemond

Data Scientist and Master Instructor
November 2015
Email: garrett@rstudio.com

Lifecycle of an Analysis Project

Clarify

Become familiar with the data,
template a solution

Develop

Create a working model

Productize

Automate and integrate

Publish

Socialize



R Markdown

Lifecycle of an Analysis Project

Clarify

Become familiar with the data,

A scripting language that is

very useful for analyzing data.

Develop

Create a working model

Productize

Automate and integrate

Publish

Socialize



R Markdown

Velocity

Volume

Data > RAM

Variety

Veracity

Lifecycle of an Analysis Project

Subset

Extract data to explore, work with

Clarify

Become familiar with the data,
template a solution

Develop

Create a working model

Productize

Automate and integrate

Publish

Socialize

*** sometimes**

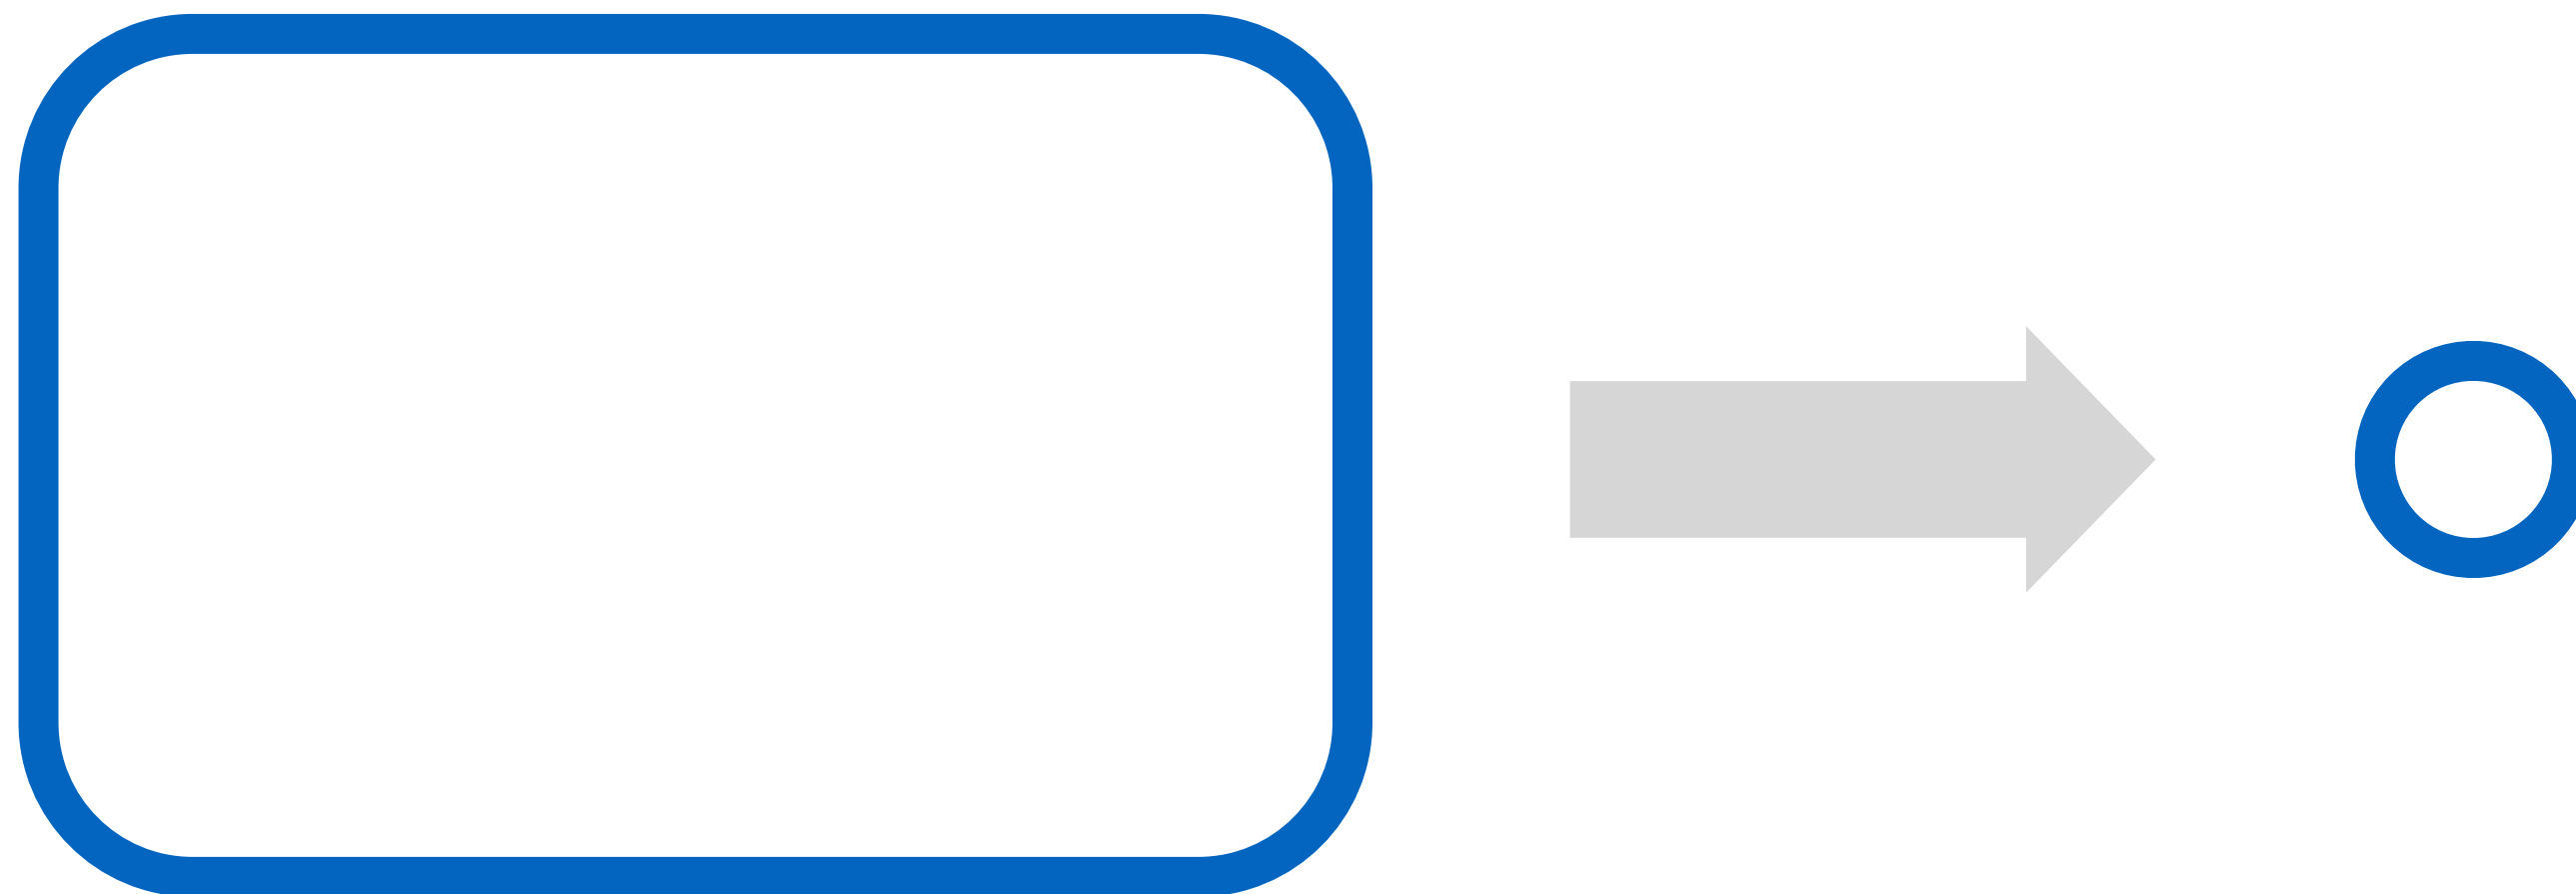
© 2015 RStudio, Inc. All rights reserved.

**How do
analysts use
big data?**

Analytic Big Data Problems

Class 1. Extract Data

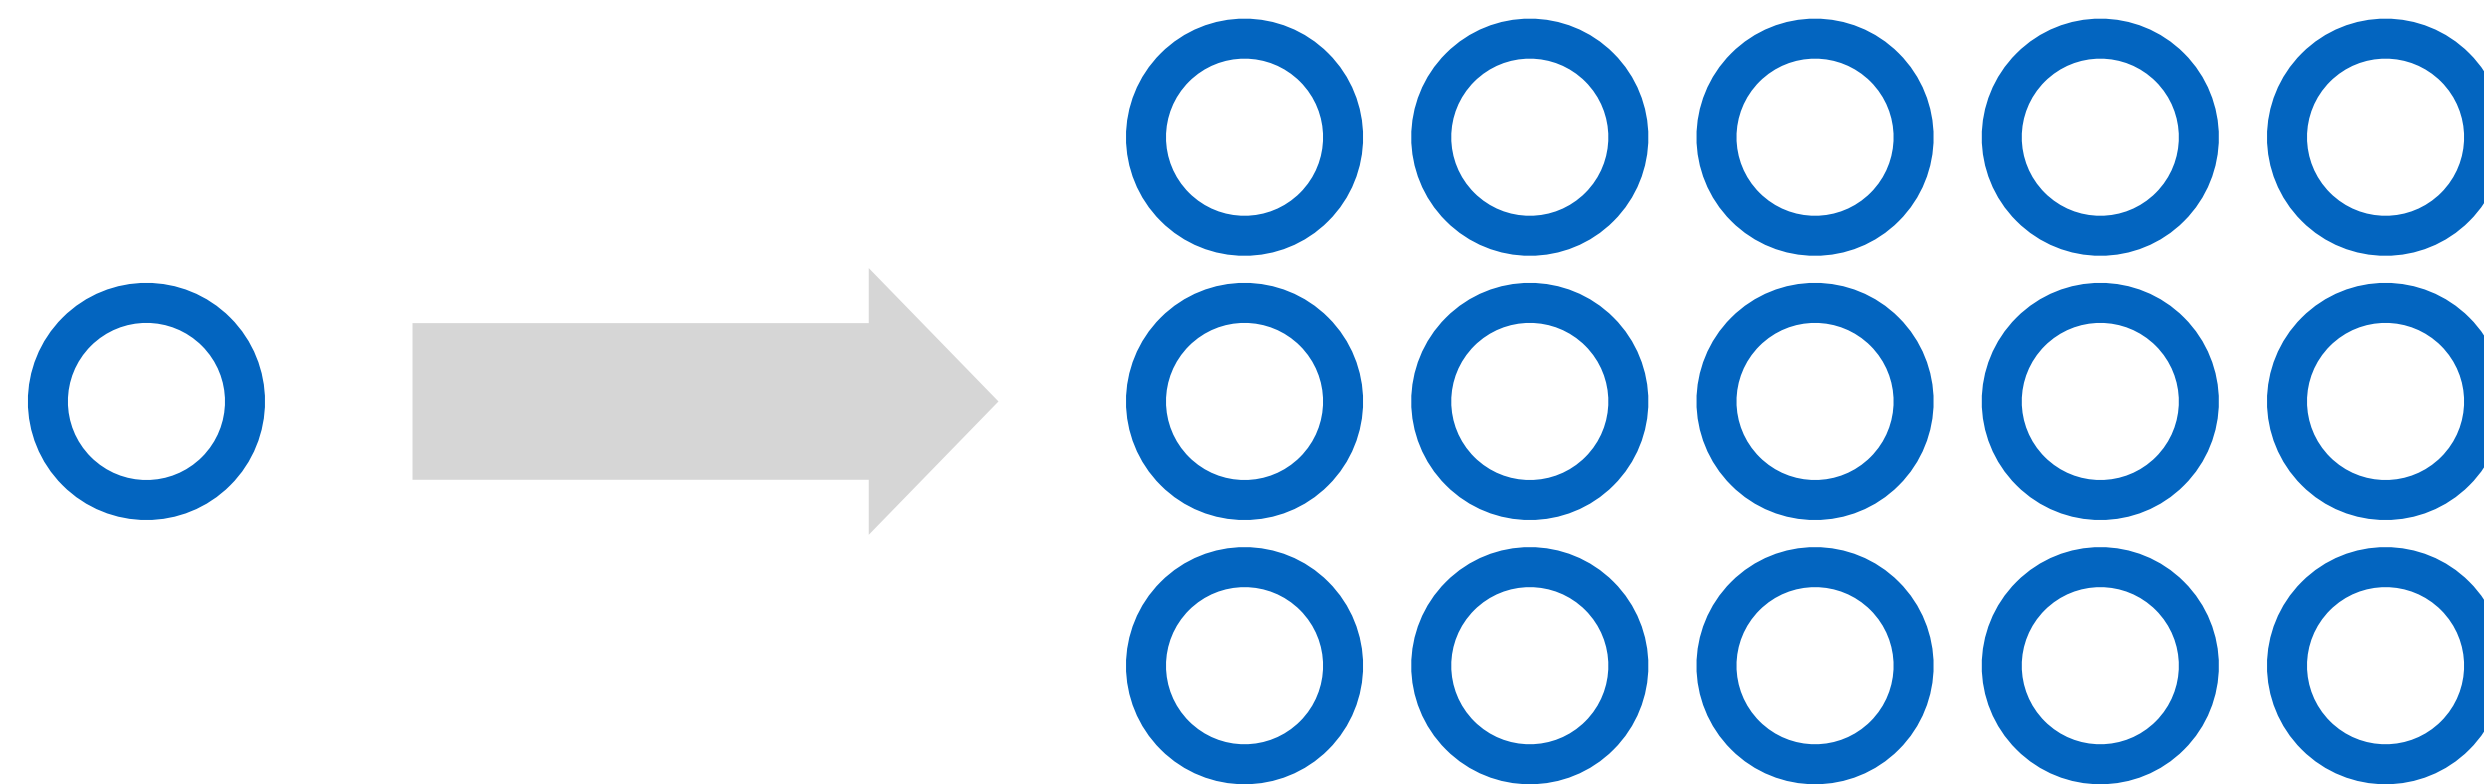
Problems that require you to extract a subset, sample, or summary from a Big Data source. You may do further analytics on the subset, and the subset might itself be quite large.



Analytic Big Data Problems

Class 2. Compute on the parts

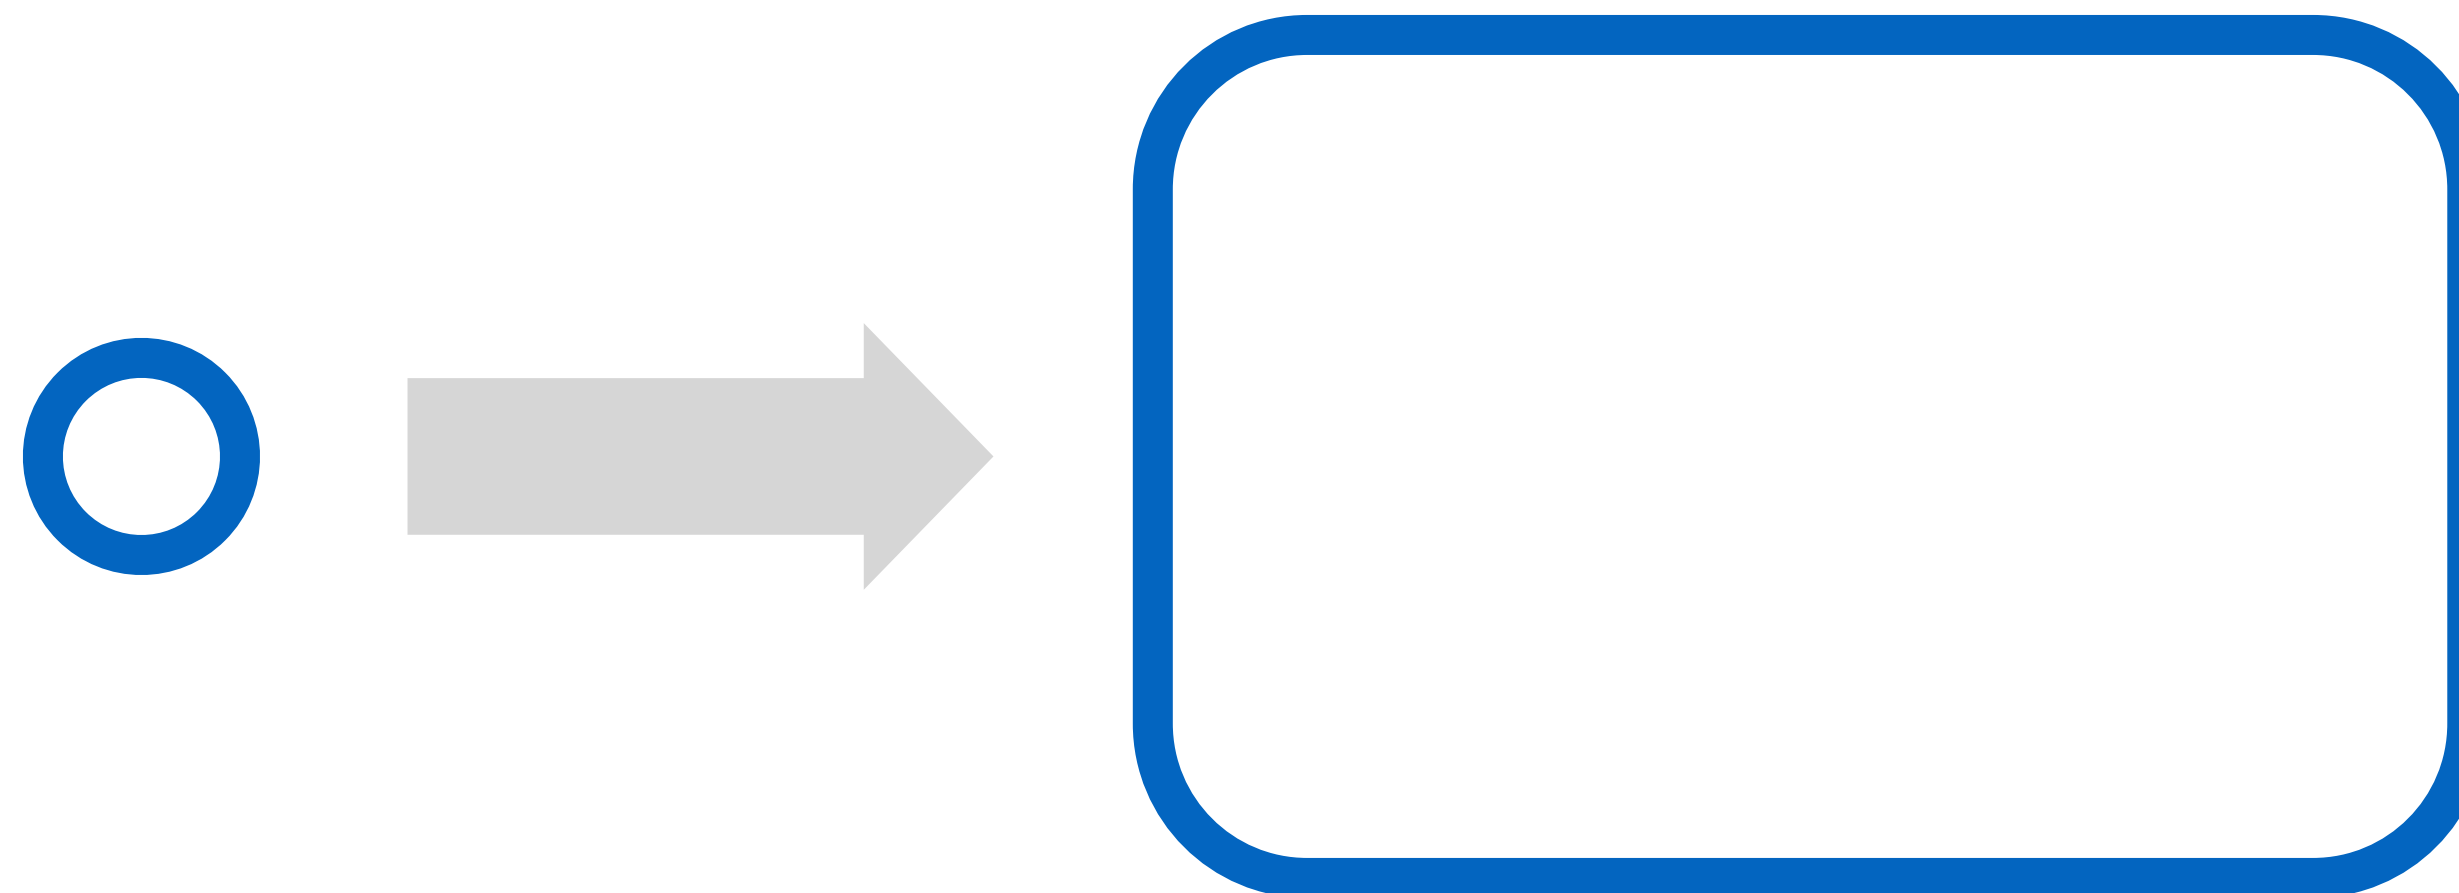
Problems that require you to repeat computation for many subgroups of the data, e.g. you need to fit one model per individual for thousands of individuals. You may combine the results once finished.



Analytic Big Data Problems

Class 3. Compute on the whole

Problems that require you to use all of the data at once. These problems are irretrievably big; they must be run at scale within the data warehouse.



Lifecycle of an Analysis Project

Subset

Extract data to explore, work with

Clarify

Become familiar with the data,
template a solution

Develop

Create a working model

Productize

Automate and integrate

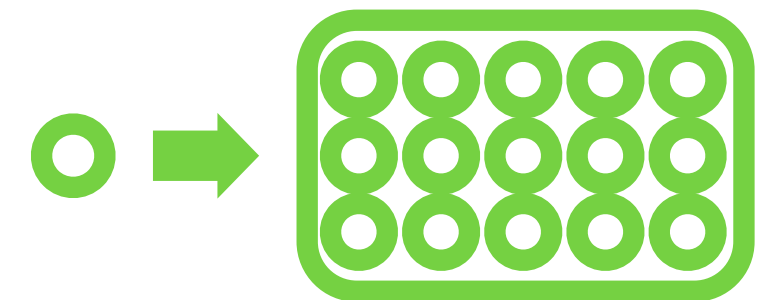
Publish

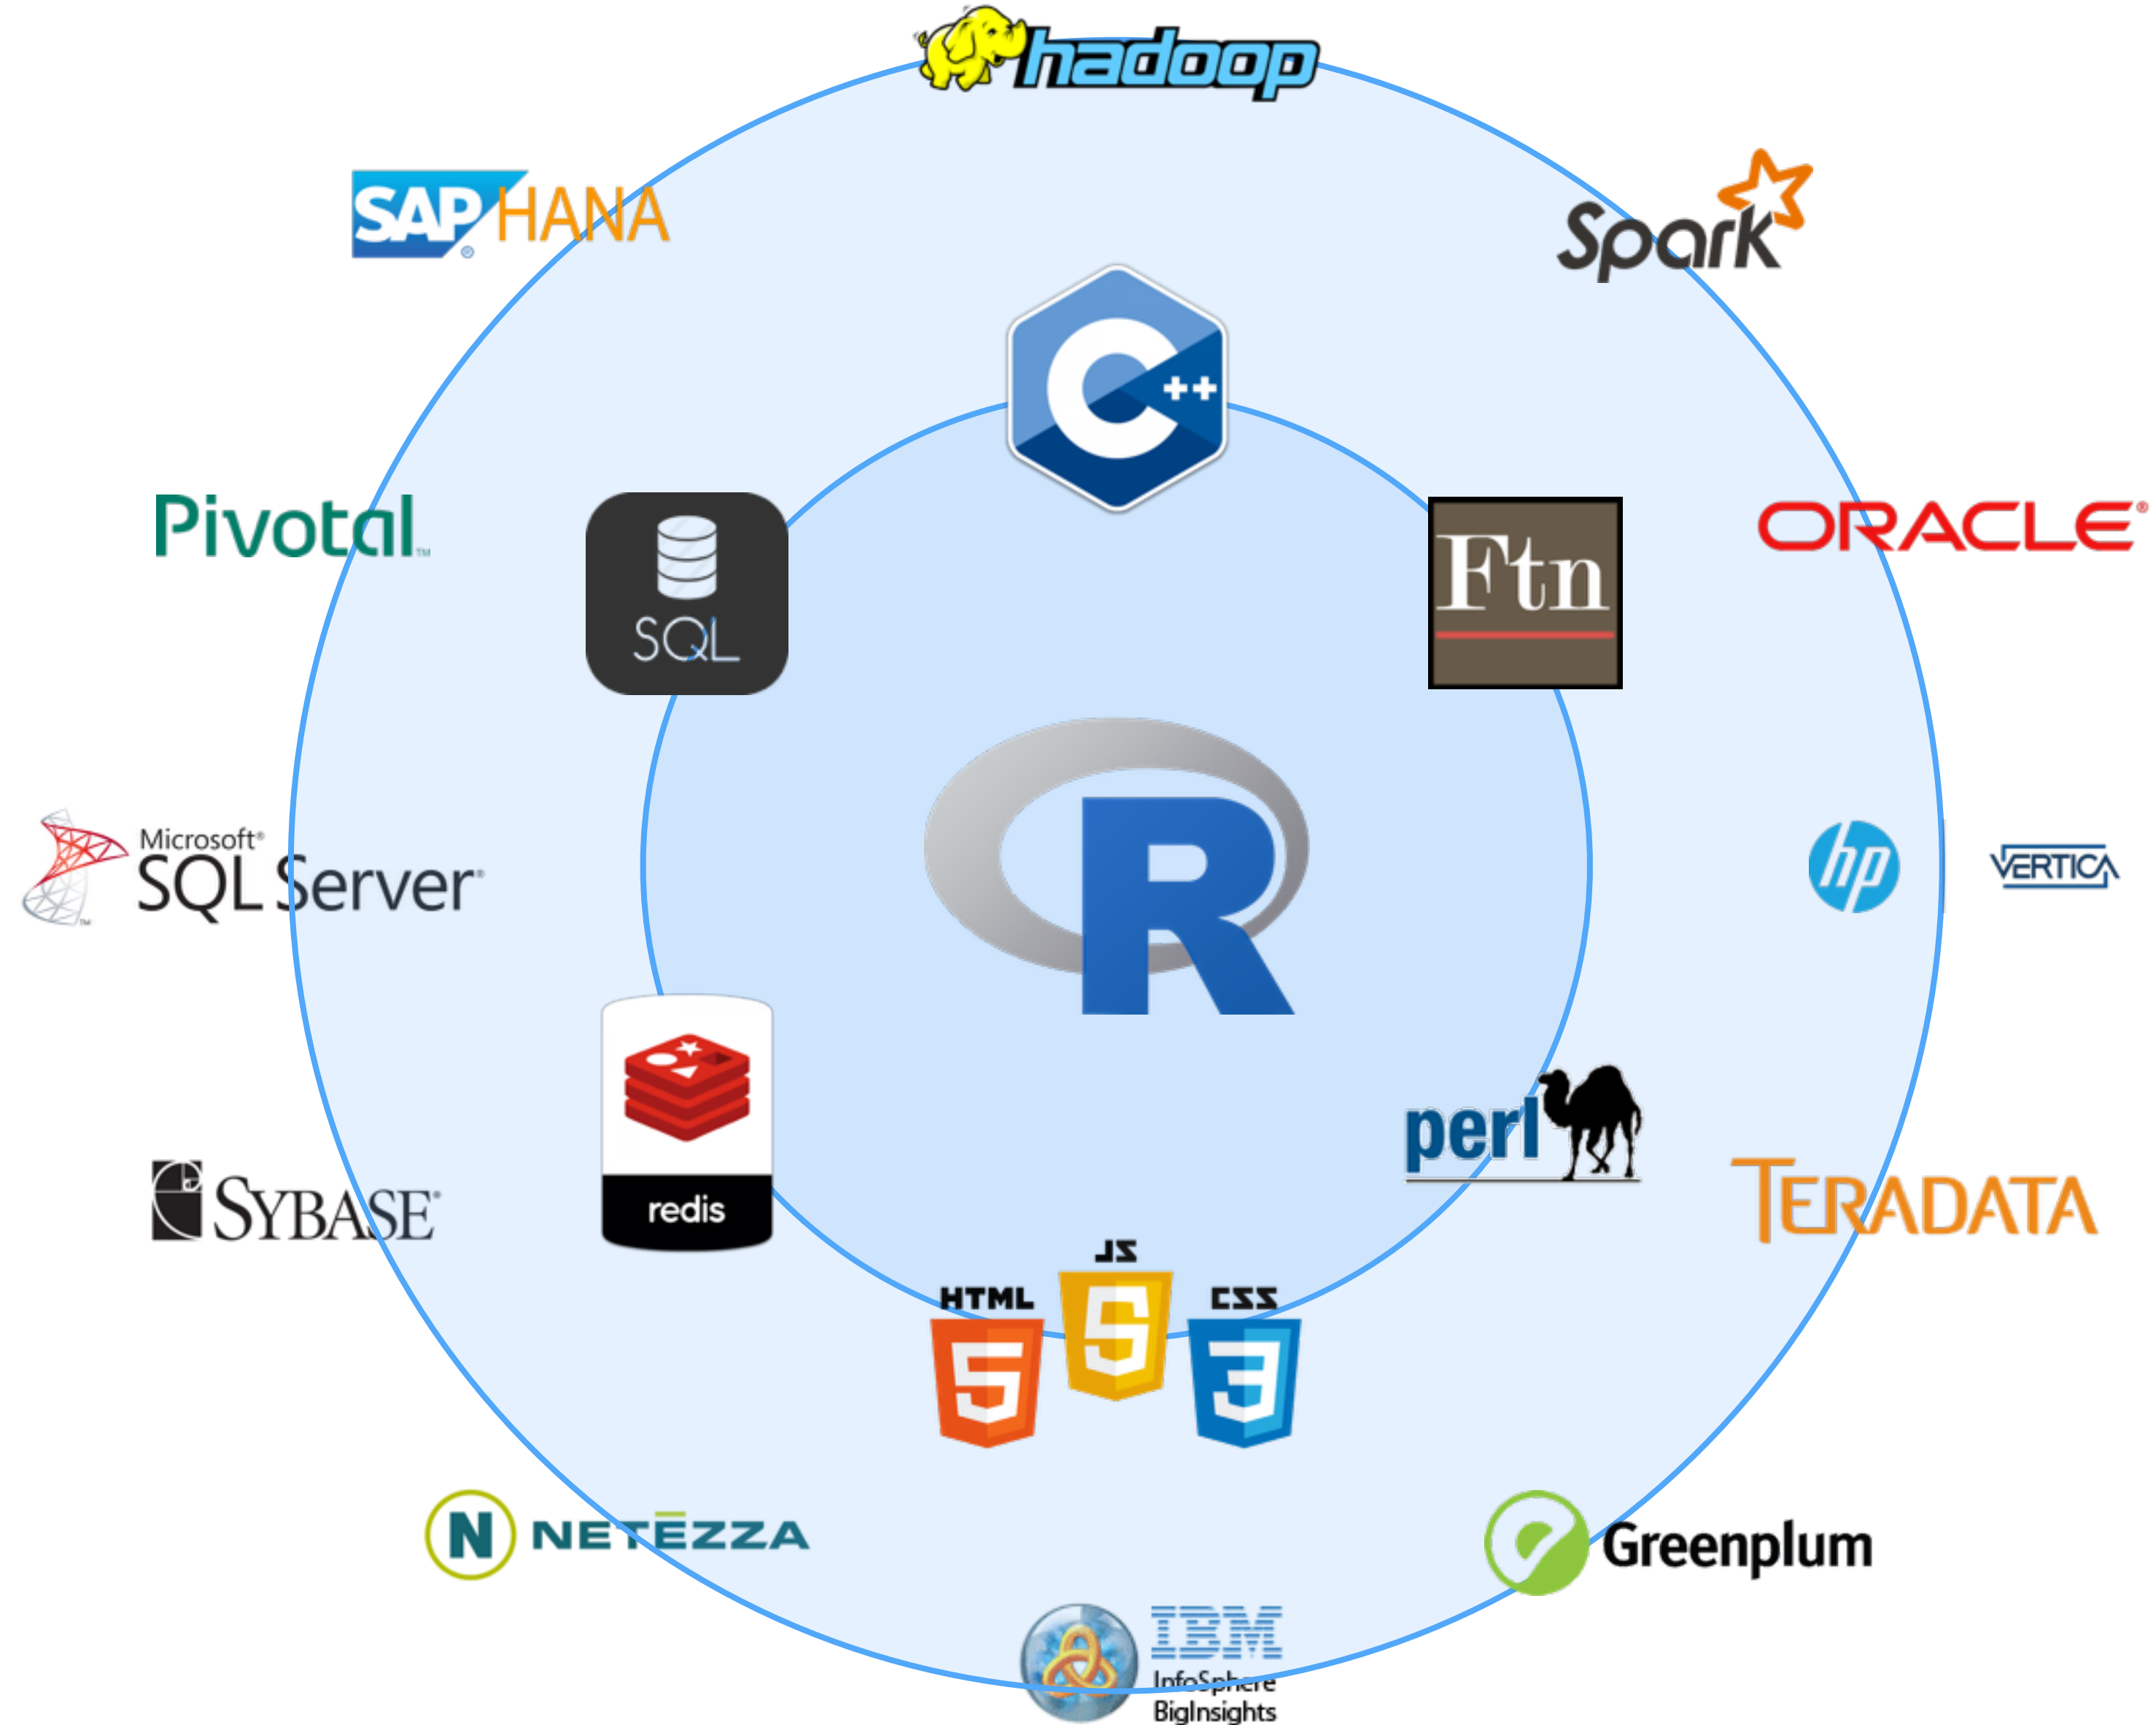
Socialize

Class 1



Class 2
Class 3

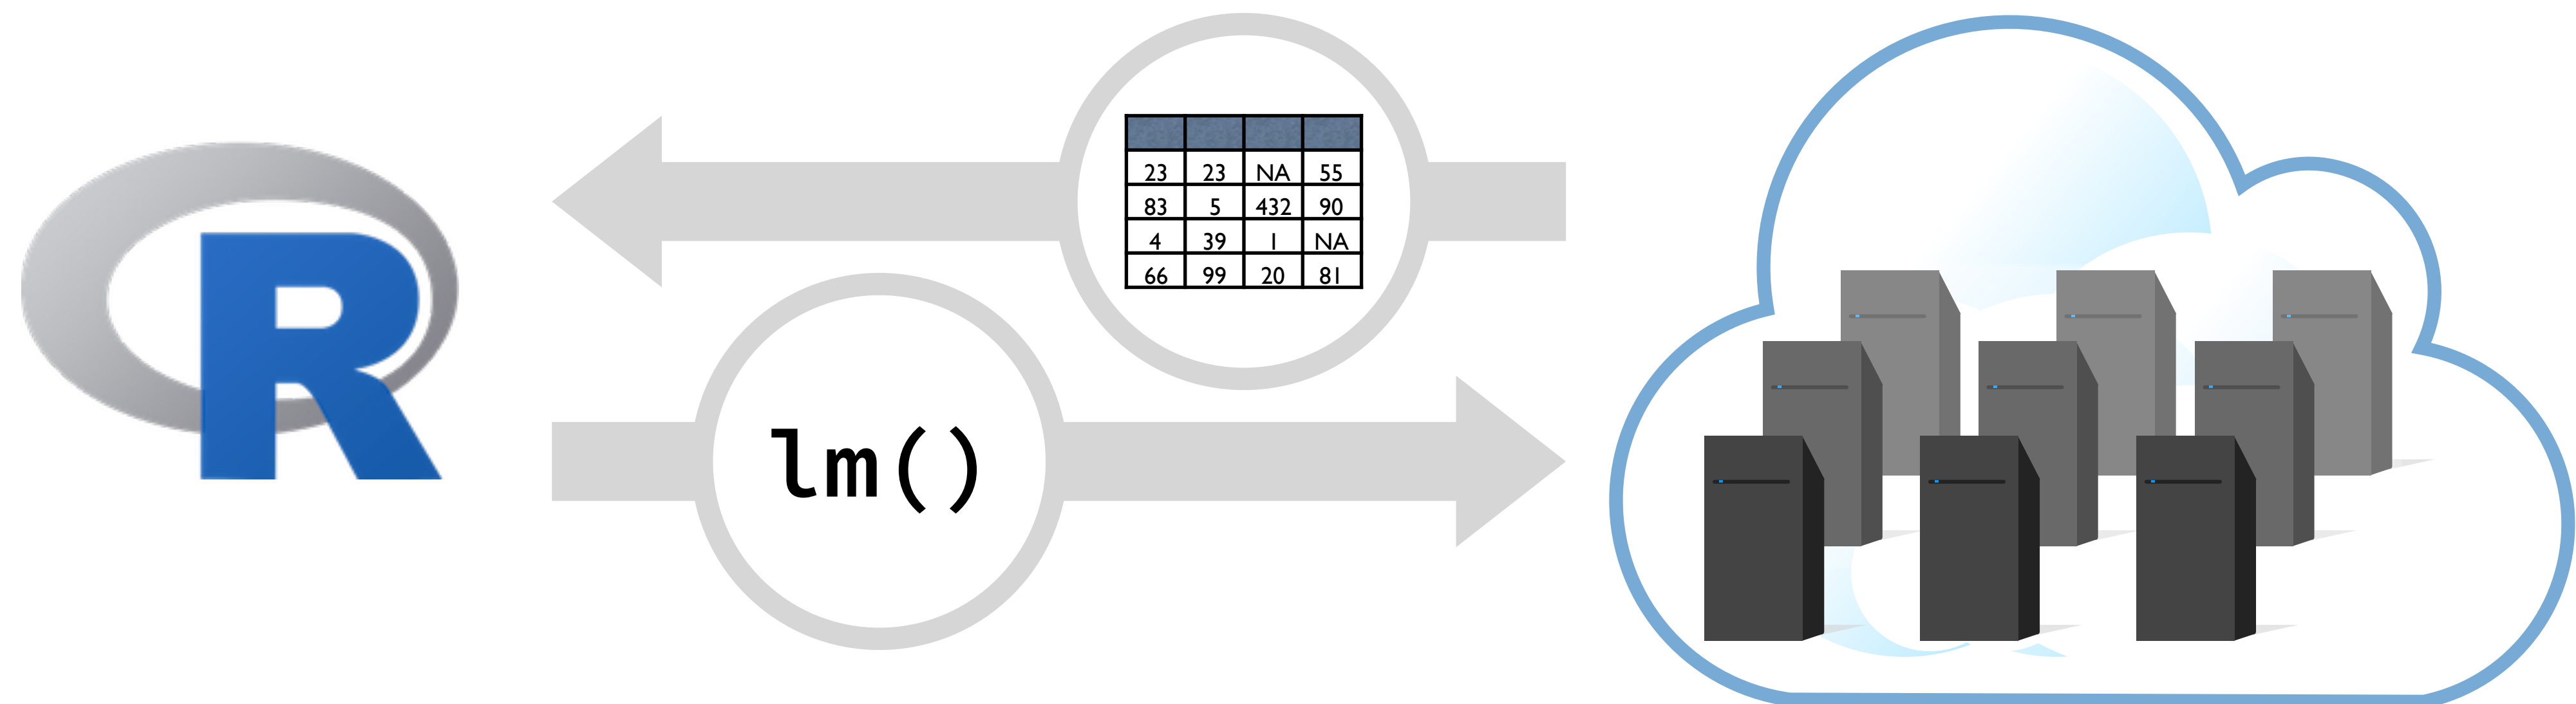




General Strategy

Store big data in a data warehouse

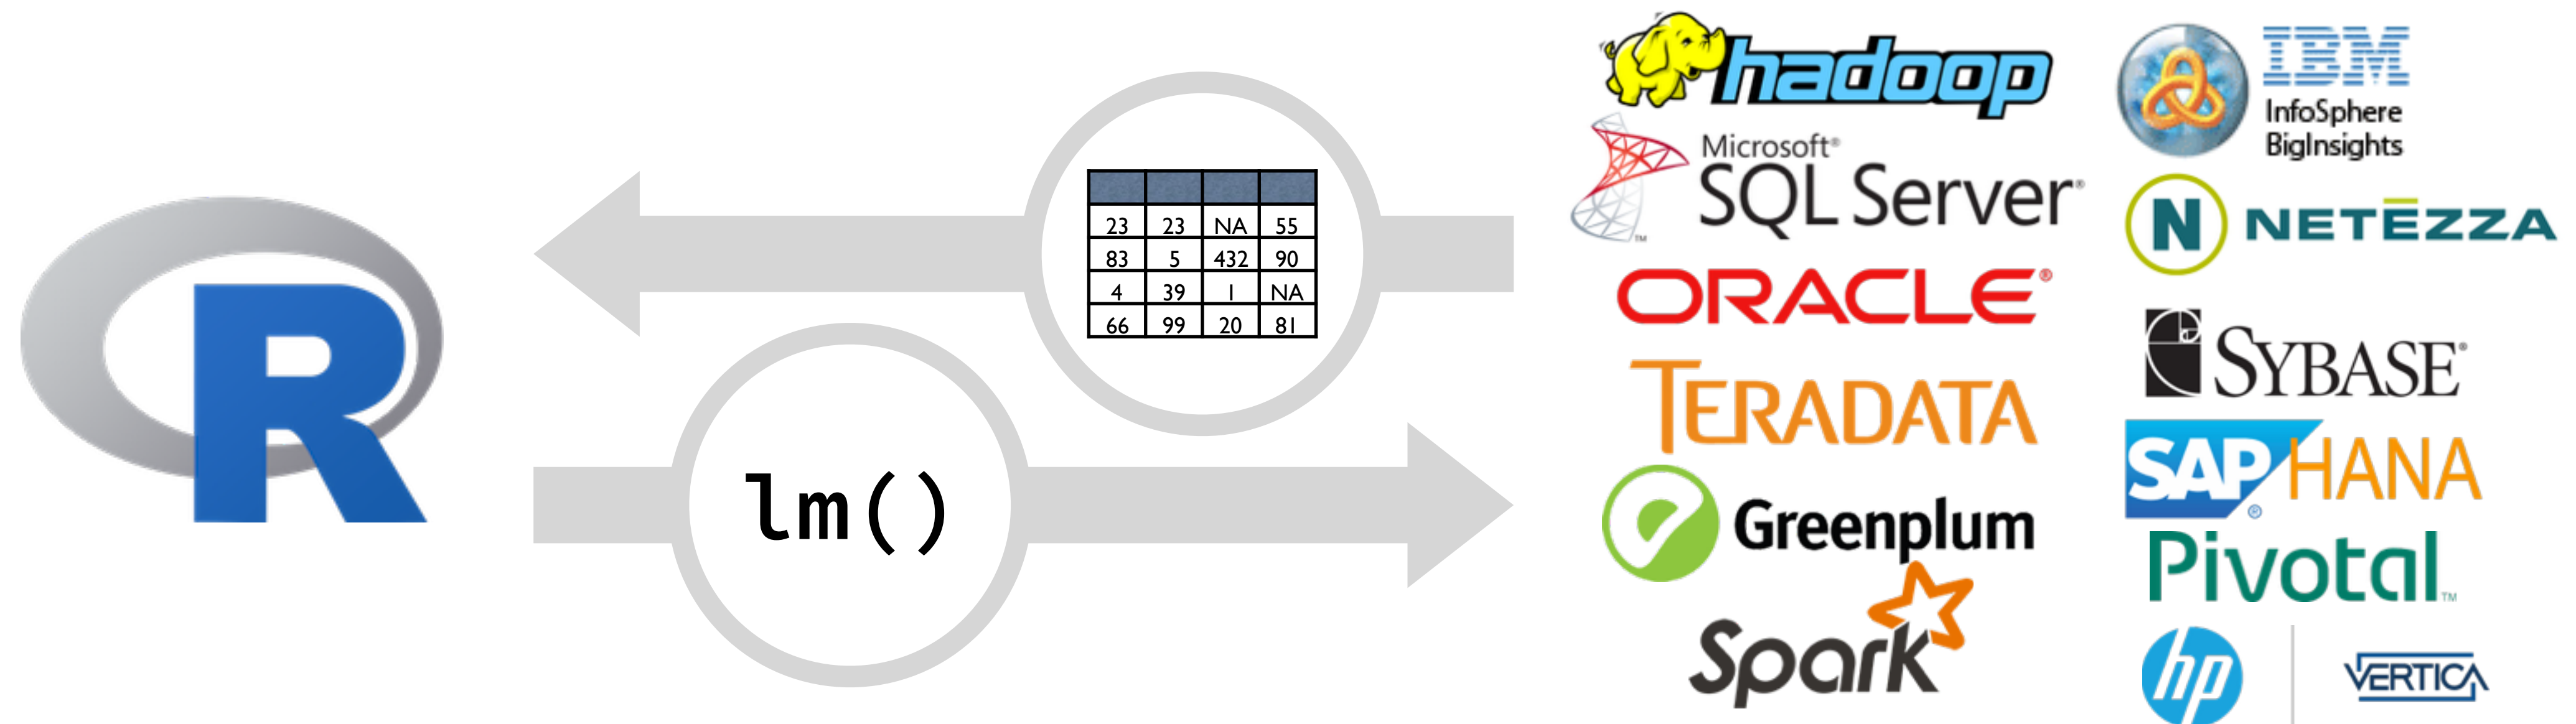
1. Pass subsets of data from warehouse to R
2. Transform R code, pass to warehouse.



General Strategy

Store big data in a data warehouse

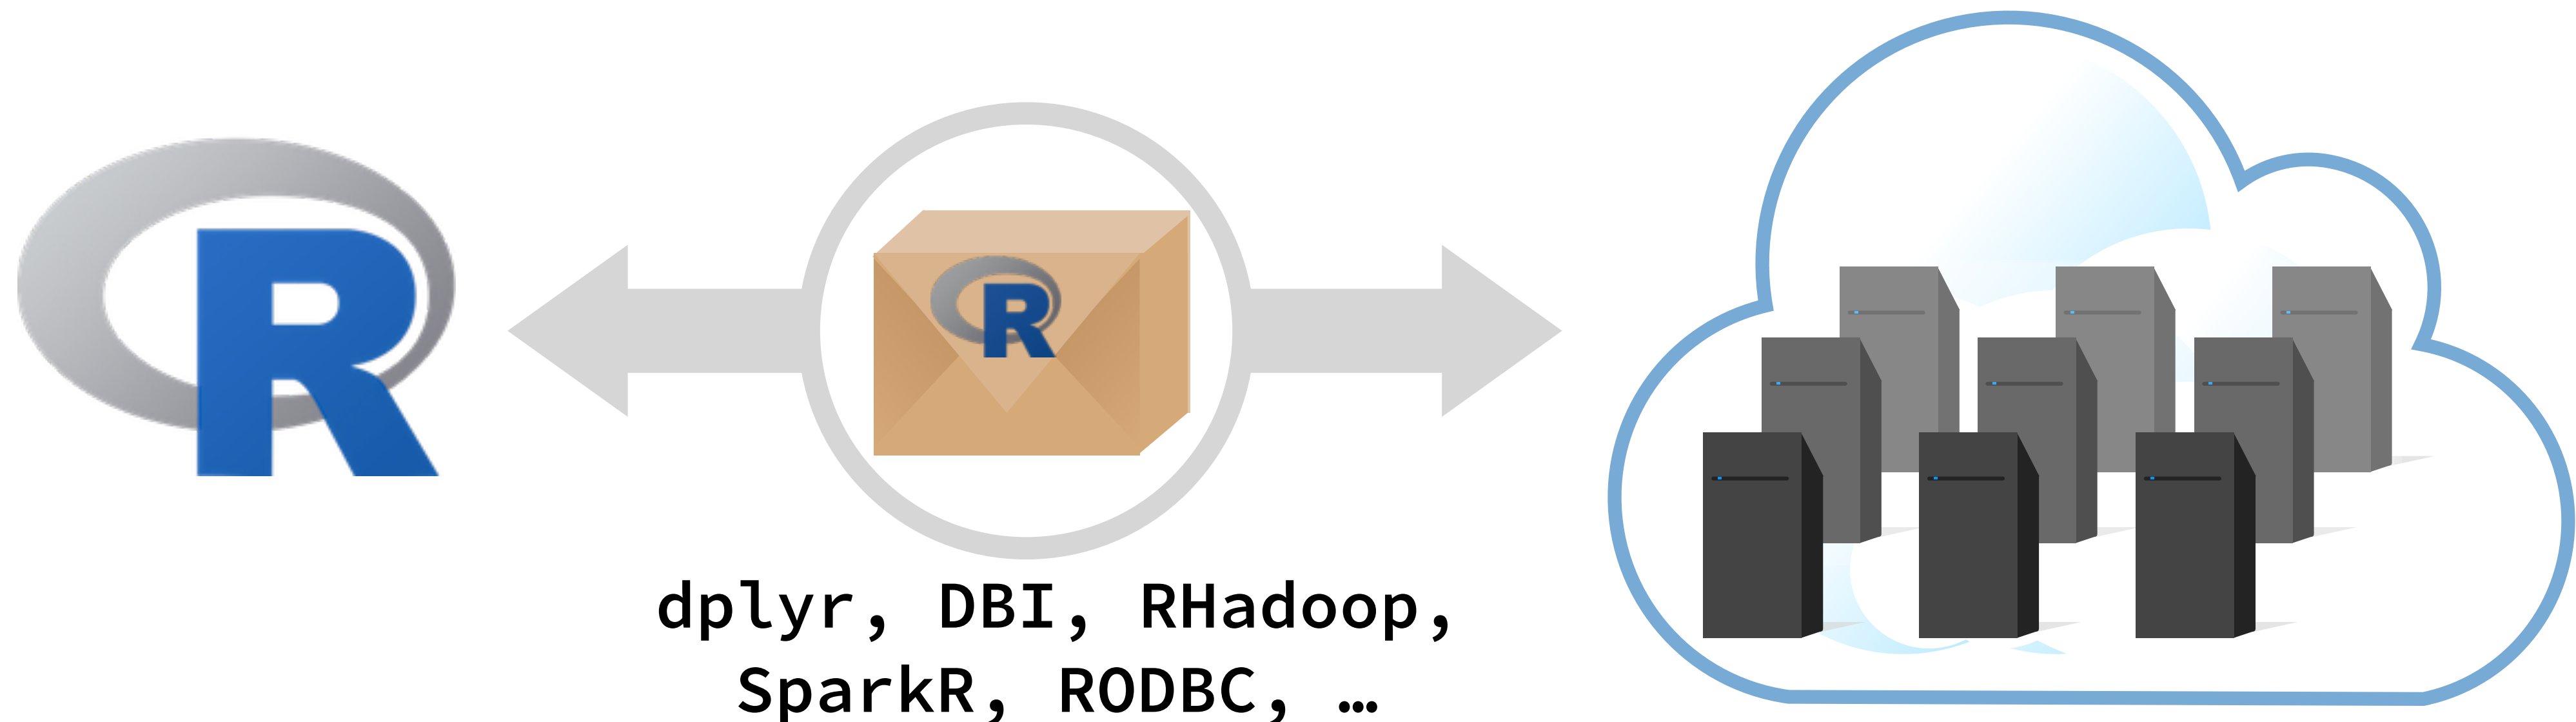
1. Pass subsets of data from warehouse to R
2. Transform R code, pass to warehouse.



General Strategy

Store big data in a data warehouse

1. Pass subsets of data from warehouse to R
2. Transform R code, pass to warehouse.



dplyr/ Redshift



✈ Airlines Data Set

Arrival and departure details for all commercial flights in US between October 1987 and April 2008.

120,000,000 records. **12 GB**

stat-computing.org/dataexpo/2009/



Data does not
fit in memory

Example Task

1. Collect random sample of training data
2. Fit a model to the sample (in R)
3. Score against test data (in DB)

```
print.R x DESCRIPTION x html_document.Rmd x handle_click.R x vega.R x
Source on Save Run Source
25 }
26
27 data_props <- combine_data_props(x$marks)
28 data_ids <- names(data_props)
29 data_table <- x$data[data_ids]
30
31 # Collapse each list of scale objects into one scale object.
32 x <- collapse_scales(x)
33 scale_data_table <- scale_domain_data(x)
34
35 # Wrap each of the reactive data objects in another reactive which returns
36 # only the columns that are actually used, and adds any calculated columns
37 # that are used in the props.
38 data_table <- active_props(data_table, data_props)
39
40 # From an environment containing data_table objects, get static data for the
41 # specified ids.
16:19 (Top Level) R Script
```

Environment History Build Git

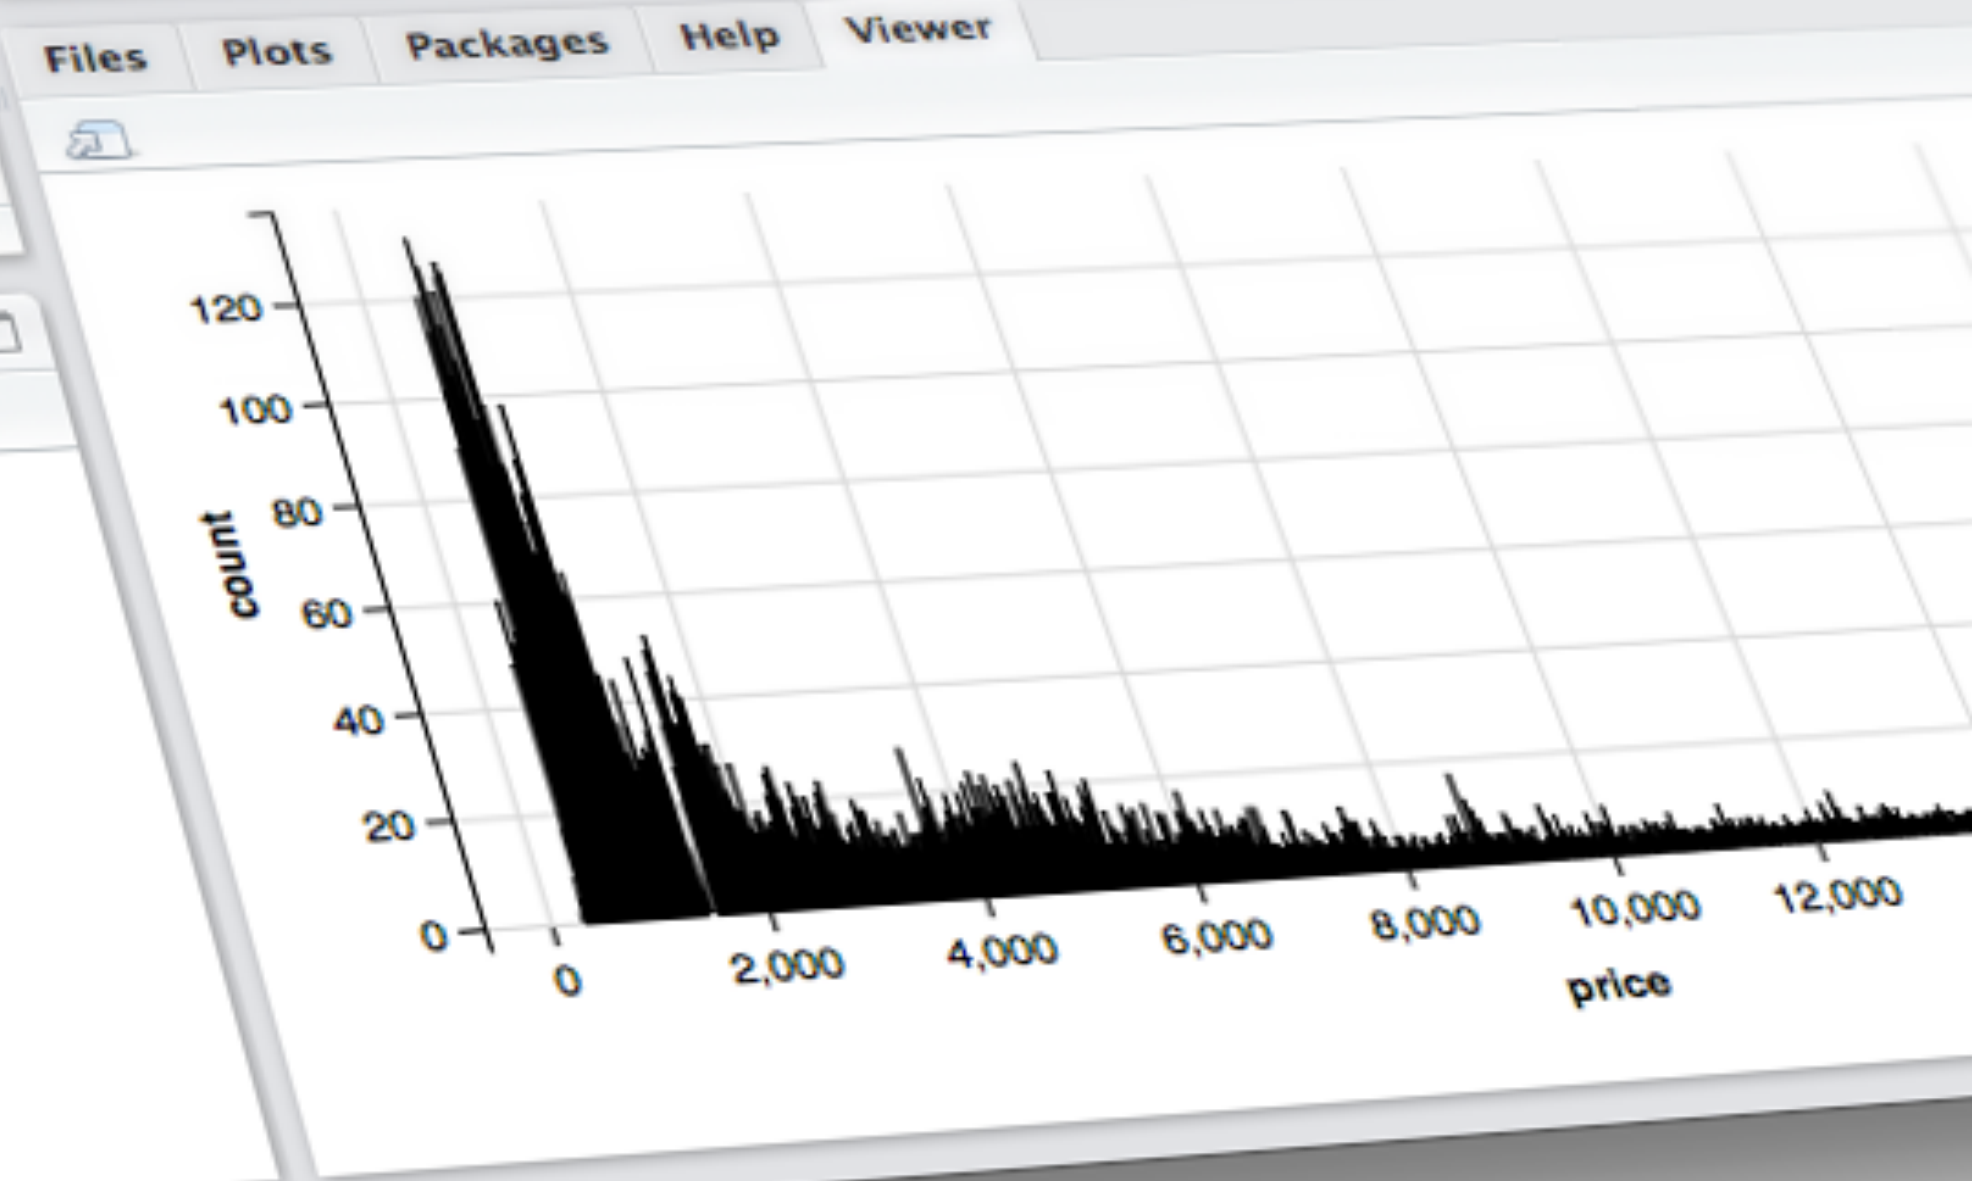
as.vega.ggvis() ~

Values

data_ids	"diamonds0/bin1/stack2"
data_props	List of 1
dynamic	FALSE

Traceback

- as.vega.ggvis(x, FALSE) at vega.R:29
- as.vega(x, FALSE) at vega.R:11
- view_static(x, ...) at print.R:67
- print.ggvis(c("list()", "list(diamonds0 = function () \nstatic_data)", "li



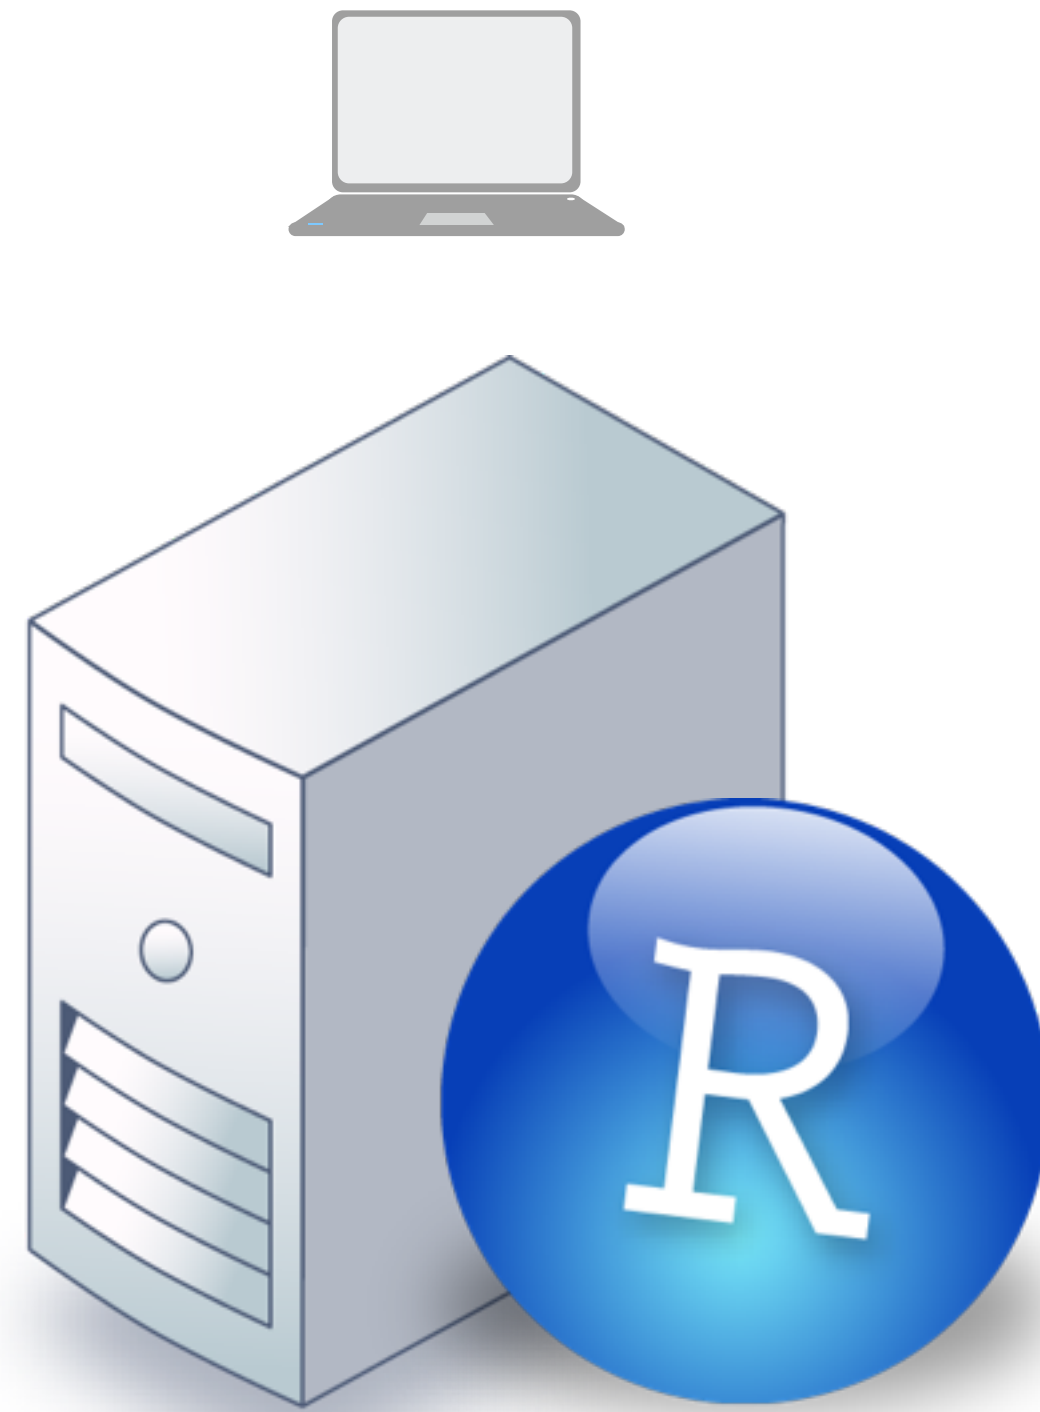
Console ~/r/ggvis/

Next Continue Stop

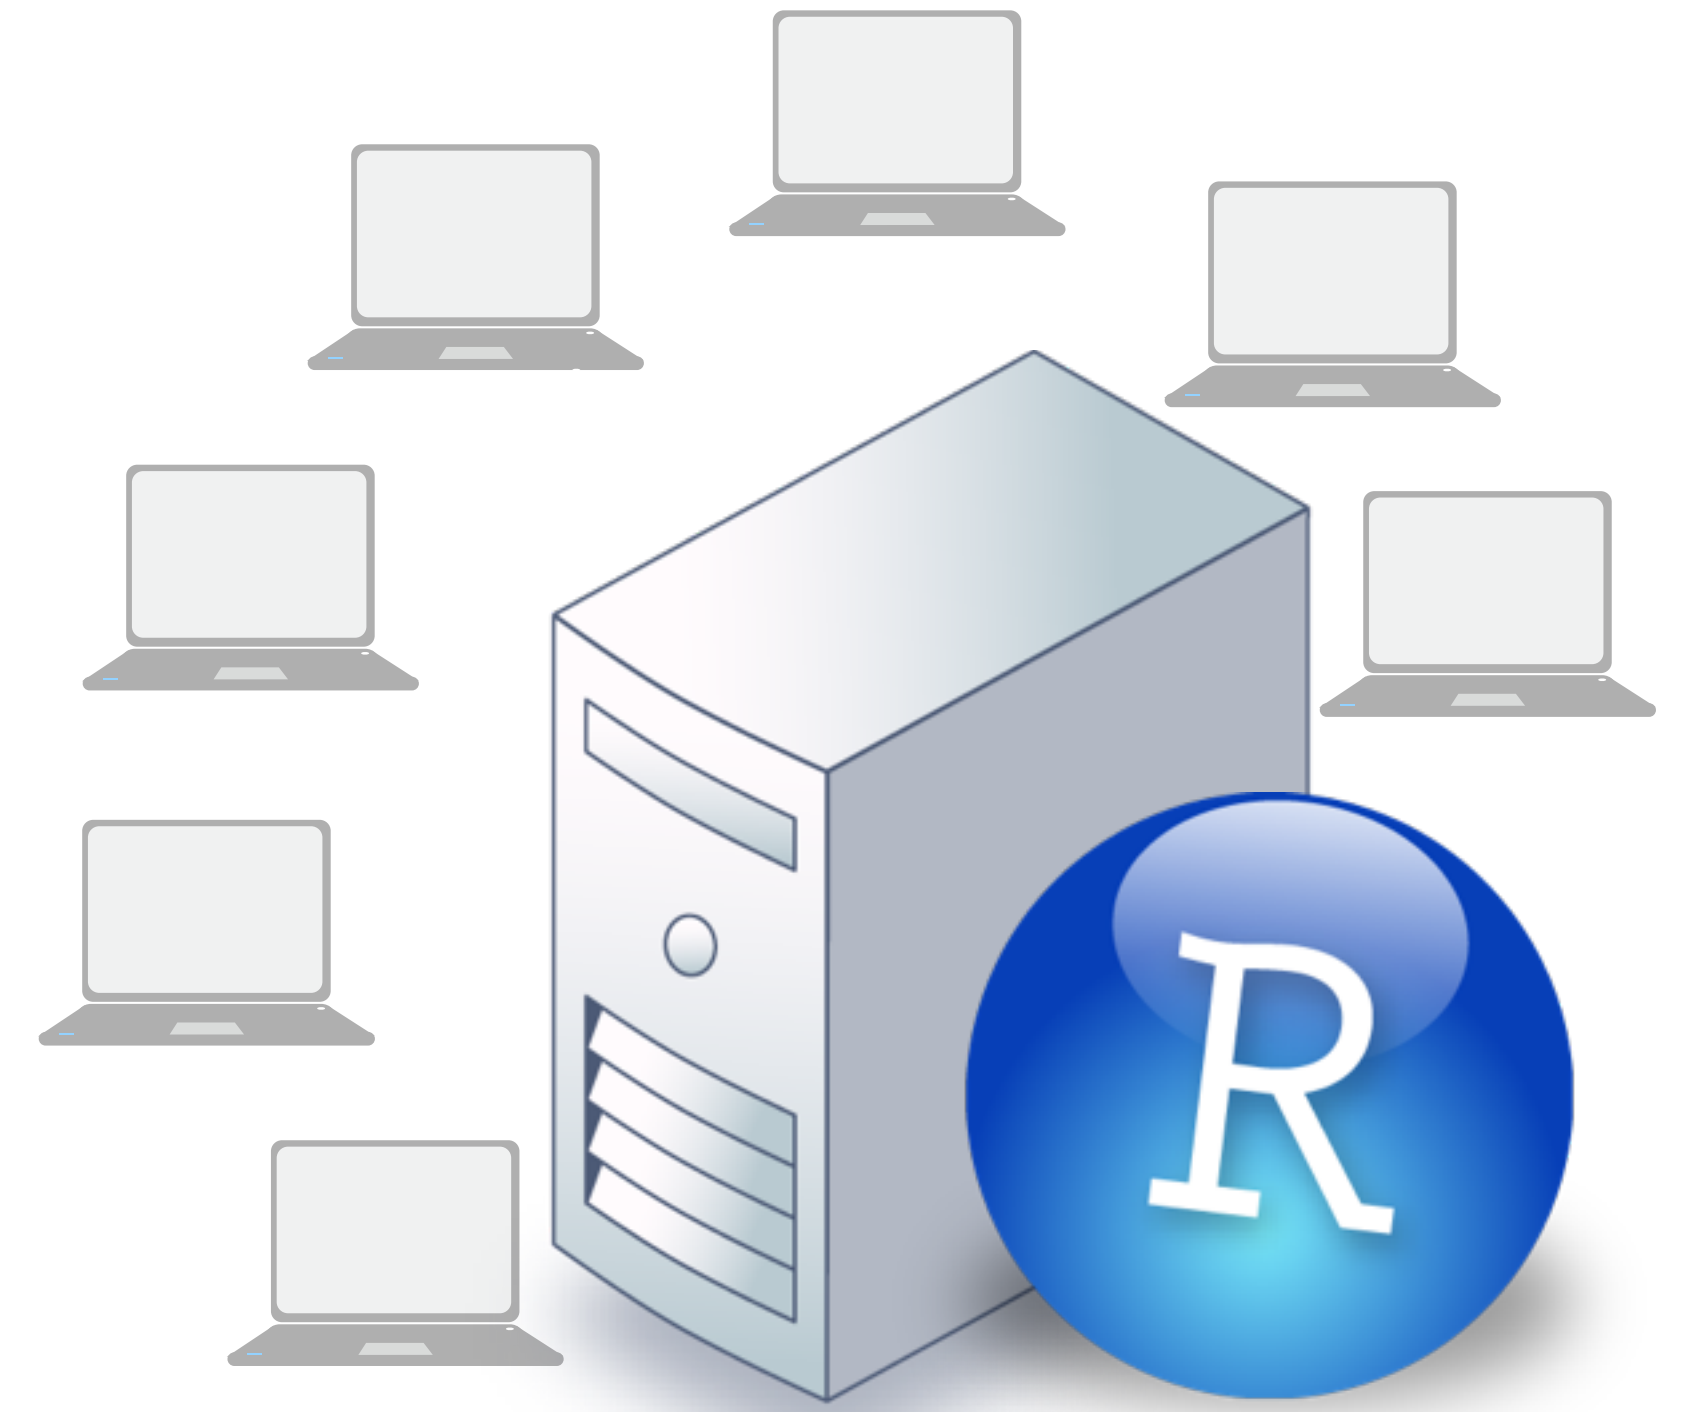
```
> ggvis(diamonds, x = ~price, y = ~color)
Guessing layer_histograms()
Guessing binwidth = 1
Called from: eval(expr, envir, enclos)
Browse[1]> n
debug at /Users/jmcphers/r/ggvis/R/vega.R#29: data_table <- x$data[data_ids]
Browse[2]>
```





RStudio
Desktop IDE



RStudio Server
Open Source



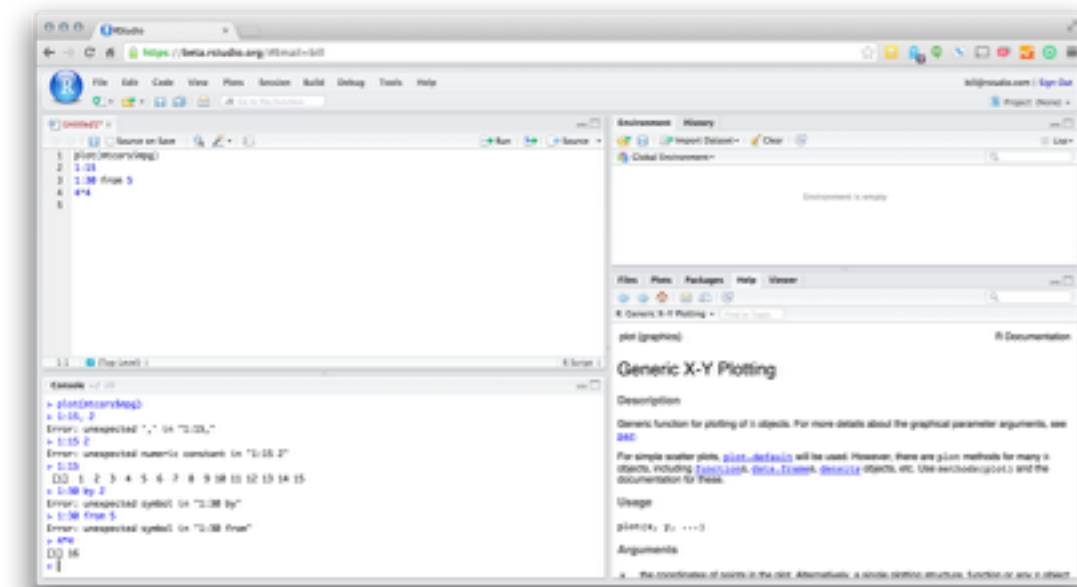
RStudio Server
Pro

- Security
 - Load balancing
 - Administrative tools
 - Resource management
 - Metrics and monitoring
- 
- RStudio
Desktop IDE

- Multiple sessions
 - Collaborative editing
 - Easy R versioning
 - Audit history
- 
- RStudio Server
Open Source



User Browser

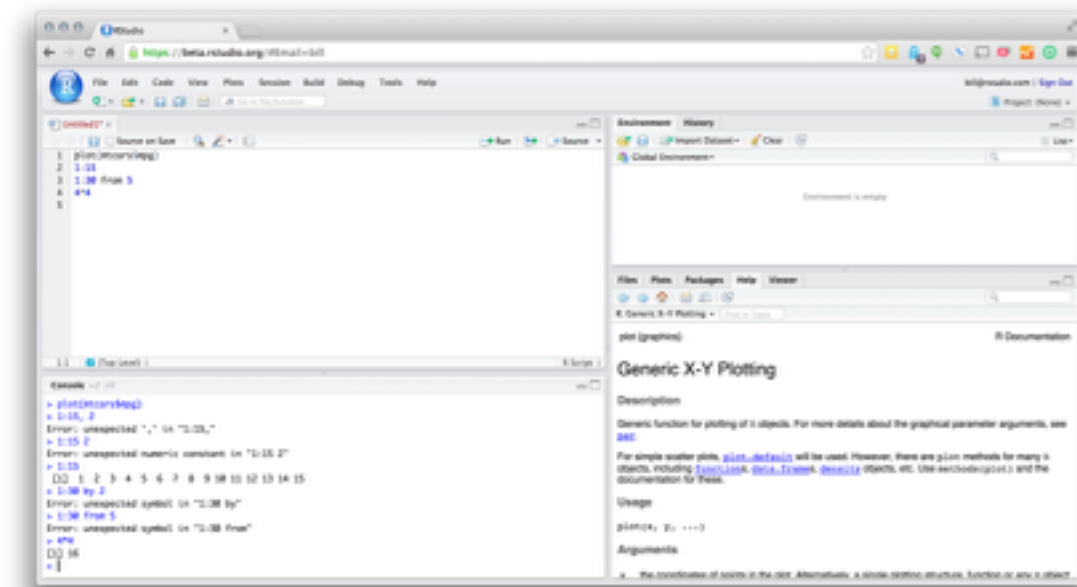


Server

RStudio Server Pro

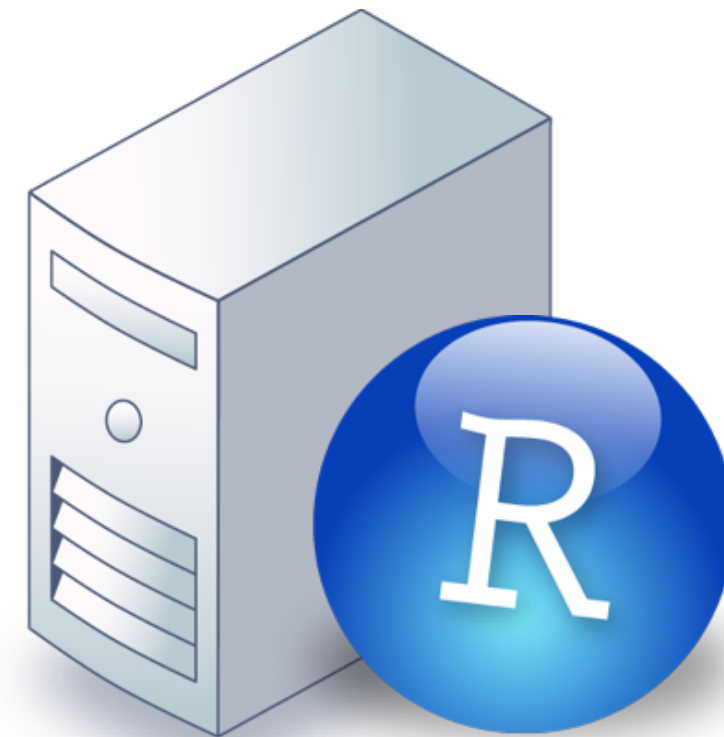


User Browser



Server

RStudio Server Pro



Database

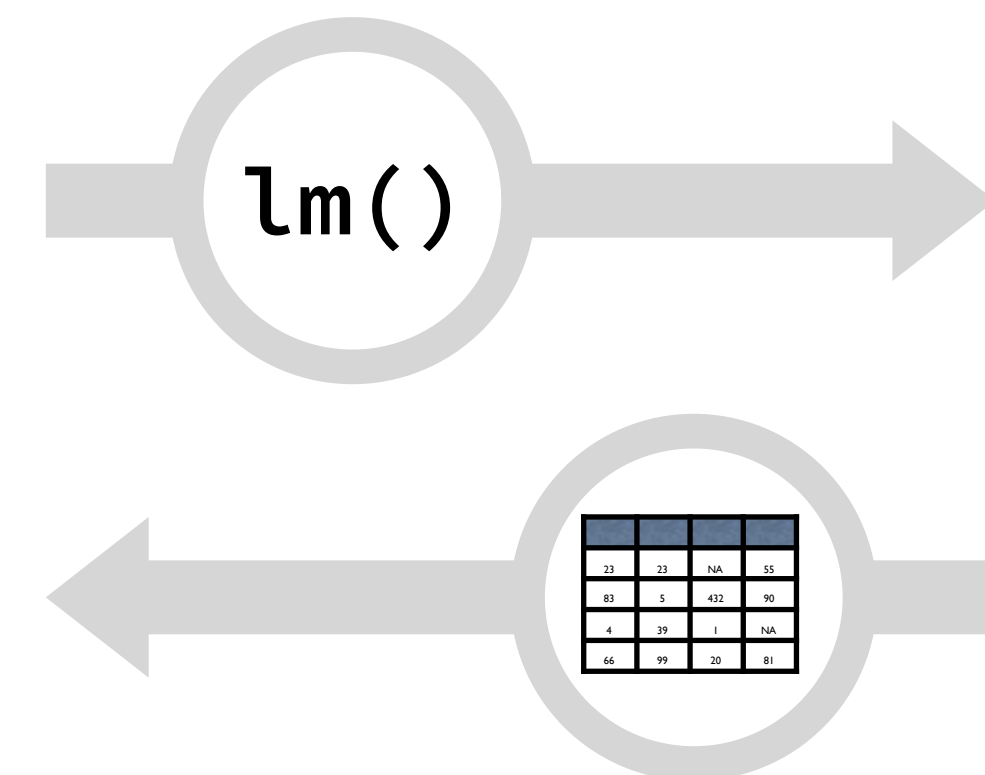
Redshift





Package that provides data manipulation syntax for R. Comes with **built-in SQL backend**:

1. **Connects** to DBMS's
2. **Transforms R code** to SQL,
sends to DBMS to run in DBMS
3. **Collect results** into R



Connections

Make with
src_function:
src_postgres, src_sqlite
src_mysql, src_bigquery

```
db <- src_postgres(  
  dbname = 'DATABASE_NAME',  
  host = 'HOST',  
  port = 5432,  
  user = 'USERNAME',  
  password = 'PASSWORD')
```

Save to
use

Driver specific
arguments.

For PostgreSQL:
dbname, host,
port, user,
password


```
library(dplyr)

# Create connection to the database
air <- src_postgres(
  dbname = 'airontime',
  host = 'sol-eng.cjku7otn8uia.us-west-2.redshift.amazonaws.com',
  port = '5439',
  user = 'redshift_user',
  password = 'ABCd4321')

# List table names
src_tbls(air)

# Create a table reference with tbl
flights <- tbl(air, "flights")
carriers <- tbl(air, "carriers")
```

Dplyr driver functions

Package	DBMS
<code>src_sqlite()</code>	SQLite
<code>src_mysql()</code>	MySQL, MariaDB
<code>src_postgres()</code>	PostgreSQL
<code>library(bigrquery)</code> <code>src_bigquery()</code>	Google BigQuery

<https://cran.r-project.org/web/packages/dplyr/vignettes/databases.html>

Adding a new SQL backend

Open guide with

```
vignette("new-sql-backend", package = "dplyr")
```

Adding a new SQL backend

2015-06-15

This document describes how to describe a new SQL backend to dplyr. It's a work in progress, but will hopefully get started. If you're familiar with how your database works, and at least one other database that dplyr already supports, this should be reasonably simple, and I'm happy to help if you get stuck. It's also possible that a new database backend may need new methods - I'm also happy to add those as needed. If you need more help, please email the [mailing list](#) and I'll help you get

```
# Manipulate the reference as if it were the actual table
clean <- flights %>%
  filter(!is.na(arrdelay), !is.na(depdelay)) %>%
  filter(depdelay > 15, depdelay < 240) %>%
  filter(year >= 2002 & year <= 2007) %>%
  select(year, arrdelay, depdelay, distance, uniquecarrier)
```

The pipe operator

%>%

```
dd <- 1:1000  
mean(dd, na.rm = TRUE)  
dd %>% mean(na.rm = TRUE)
```

These do the
same thing

Try it!



dd

mean(_____, na.rm = TRUE)

show_query

To see the SQL that dplyr will run.

```
show_query(clean)
```

show_query

tbl made from table reference
and dplyr functions

```
show_query(clean)
```

```
## <SQL>
```

```
## SELECT "year" AS "year",
```

```
##     "arrdelay" AS "arrdelay",
```

```
##     "depdelay" AS "depdelay",
```

```
##     "distance" AS "distance",
```

```
##     "uniquecarrier" AS "uniquecarrier"
```

```
## FROM "flights"
```

```
## WHERE NOT("arrdelay" IS NULL) AND NOT("depdelay" IS NULL)
```

```
##     AND "depdelay" > 15.0 AND "depdelay" < 240.0
```

```
##     AND "year" >= 2002.0 AND "year" <= 2007.0
```


** dplyr can convert all of the following to SQL*
*** other functions will be passed as is into SQL*

dplyr functions

arrange, filter, group_by, mutate, select, summarize, %>%, left_join, etc.

Operators

+, -, *, /, %%, ^

Math functions

abs, acos, cosh, sin, asinh, atan, atan2, atanh, ceiling, cos, cosh, cot, coth, exp, floor, log, log10, round, sign, sin, sinh, sqrt, tan, tanh

Comparisons

<, <=, !=, >=, >, ==, %in%

Booleans

&, &&, |, ||, !, xor

Aggregations

mean, sum, min, max, sd, var

Lazy Execution 1

```
q1 <- filter(flights, year < 2007)
q2 <- filter(q1, depdelay > 15)
q3 <- filter(q2, depdelay < 240)
q4 <- select(q3, arrdelay, depdelay, year)
q4
```

When should dplyr query
the database?

Lazy Execution 1

```
q1 <- filter(flights, year < 2007)
q2 <- filter(q1, depdelay > 15)
q3 <- filter(q2, depdelay < 240)
q4 <- select(q3, arrdelay, depdelay, year)
q4
```

dplyr will not retrieve data until last possible moment. It combines all necessary work into a single query.

```
show_query(q4)
```

```
## <SQL>
```

```
## SELECT "arrdelay" AS "arrdelay",  
         "depdelay" AS "depdelay",  
         "year" AS "year"
```

```
## FROM "flights"
```

```
## WHERE "year" > 2007.0  
        AND "depdelay" > 15.0  
        AND "depdelay" < 240.0
```

Lazy Execution 2

dplyr will only retrieve the **first 10 rows** of a query when you look at the output.

```
clean
## Source: postgres 8.0.2 [...]
## From: flights [6,517,621 x 4]
## Filter: !is.na(arrdelay), !is.na(depdelay), ...

##   year arrdelay depdelay uniquecarrier
##   (int)   (int)   (int)         (chr)
## 1  2007     42     40             9E
## 2  2007     90     94             9E
## 3  2007     19     20             9E
## 4  2007    184    167             9E
## 5  2007     21     30             9E
## 6  2007    178    179             9E
## 7  2007     56     59             9E
## 8  2007     21     21             9E
## 9  2007     50     57             9E
## 10 2007     56     23             9E
## .. ... .. ... ..
```

collect()

Forces dplyr to retrieve all results into **R**.

```
q5 <- flights %>%  
  filter(year > 2007, depdelay > 15) %>%  
  filter(depdelay == 240) %>%  
  collect()
```

collect() returns entire
result as a tbl_df

collect()

query to run

collapse()

Forces execution in **DBMS**

```
q6 <- flights %>%  
  mutate(adjdelay = depdelay - 15) %>%  
  collapse() %>%  
  filter(adjdelay > 0)
```

collapse() turns the preceding queries into a table expression

remaining queries are run against the table described in the collapsed expression

```
# Extract random 1% sample of training data
random <- clean %>%
  mutate(x = random()) %>%
  collapse() %>%
  filter(x <= 0.01) %>%
  select(-x) %>%
  collect()
```

Fit a model (in R)

Do some carriers make up lost time better than others?

$$gain = depdelay - arrdelay$$


```
# make gain
random$gain <- random$depdelay - random$arrdelay

# build model
mod <- lm(gain ~ depdelay + distance + uniquecarrier,
          data = random)

# carrier coefficients table
coef(mod)
```

```
# make coefficients lookup table
coefs <- dummy.coef(mod)

coefs_table <- data.frame(
  uniquecarrier = names(coefs$uniquecarrier),
  carrier_score = coefs$uniquecarrier,
  int_score = coefs$`(Intercept)`,
  dist_score = coefs$distance,
  delay_score = coefs$depdelay,
  row.names = NULL,
  stringsAsFactors = FALSE
)
```

coefs_table

##	uniquecarrier	carrier_score	int_score	dist_score	delay_score
## 1	9E	0.0000000	-1.540312	0.003083624	-0.01359926
## 2	AA	-1.7131012	-1.540312	0.003083624	-0.01359926
## 3	AQ	0.6153050	-1.540312	0.003083624	-0.01359926
## 4	AS	1.4143664	-1.540312	0.003083624	-0.01359926
## 5	B6	-1.9714287	-1.540312	0.003083624	-0.01359926
## 6	C0	-1.5865993	-1.540312	0.003083624	-0.01359926
## 7	DH	3.1367039	-1.540312	0.003083624	-0.01359926
## 8	DL	-2.6404154	-1.540312	0.003083624	-0.01359926
## 9	EV	2.3434536	-1.540312	0.003083624	-0.01359926
## 10	F9	0.5341914	-1.540312	0.003083624	-0.01359926
## 11	FL	-0.8888280	-1.540312	0.003083624	-0.01359926
## 12	HA	1.6712540	-1.540312	0.003083624	-0.01359926
## 13	HP	3.3742529	-1.540312	0.003083624	-0.01359926
## 14	MQ	-1.3632398	-1.540312	0.003083624	-0.01359926
## 15	NW	-2.0416490	-1.540312	0.003083624	-0.01359926

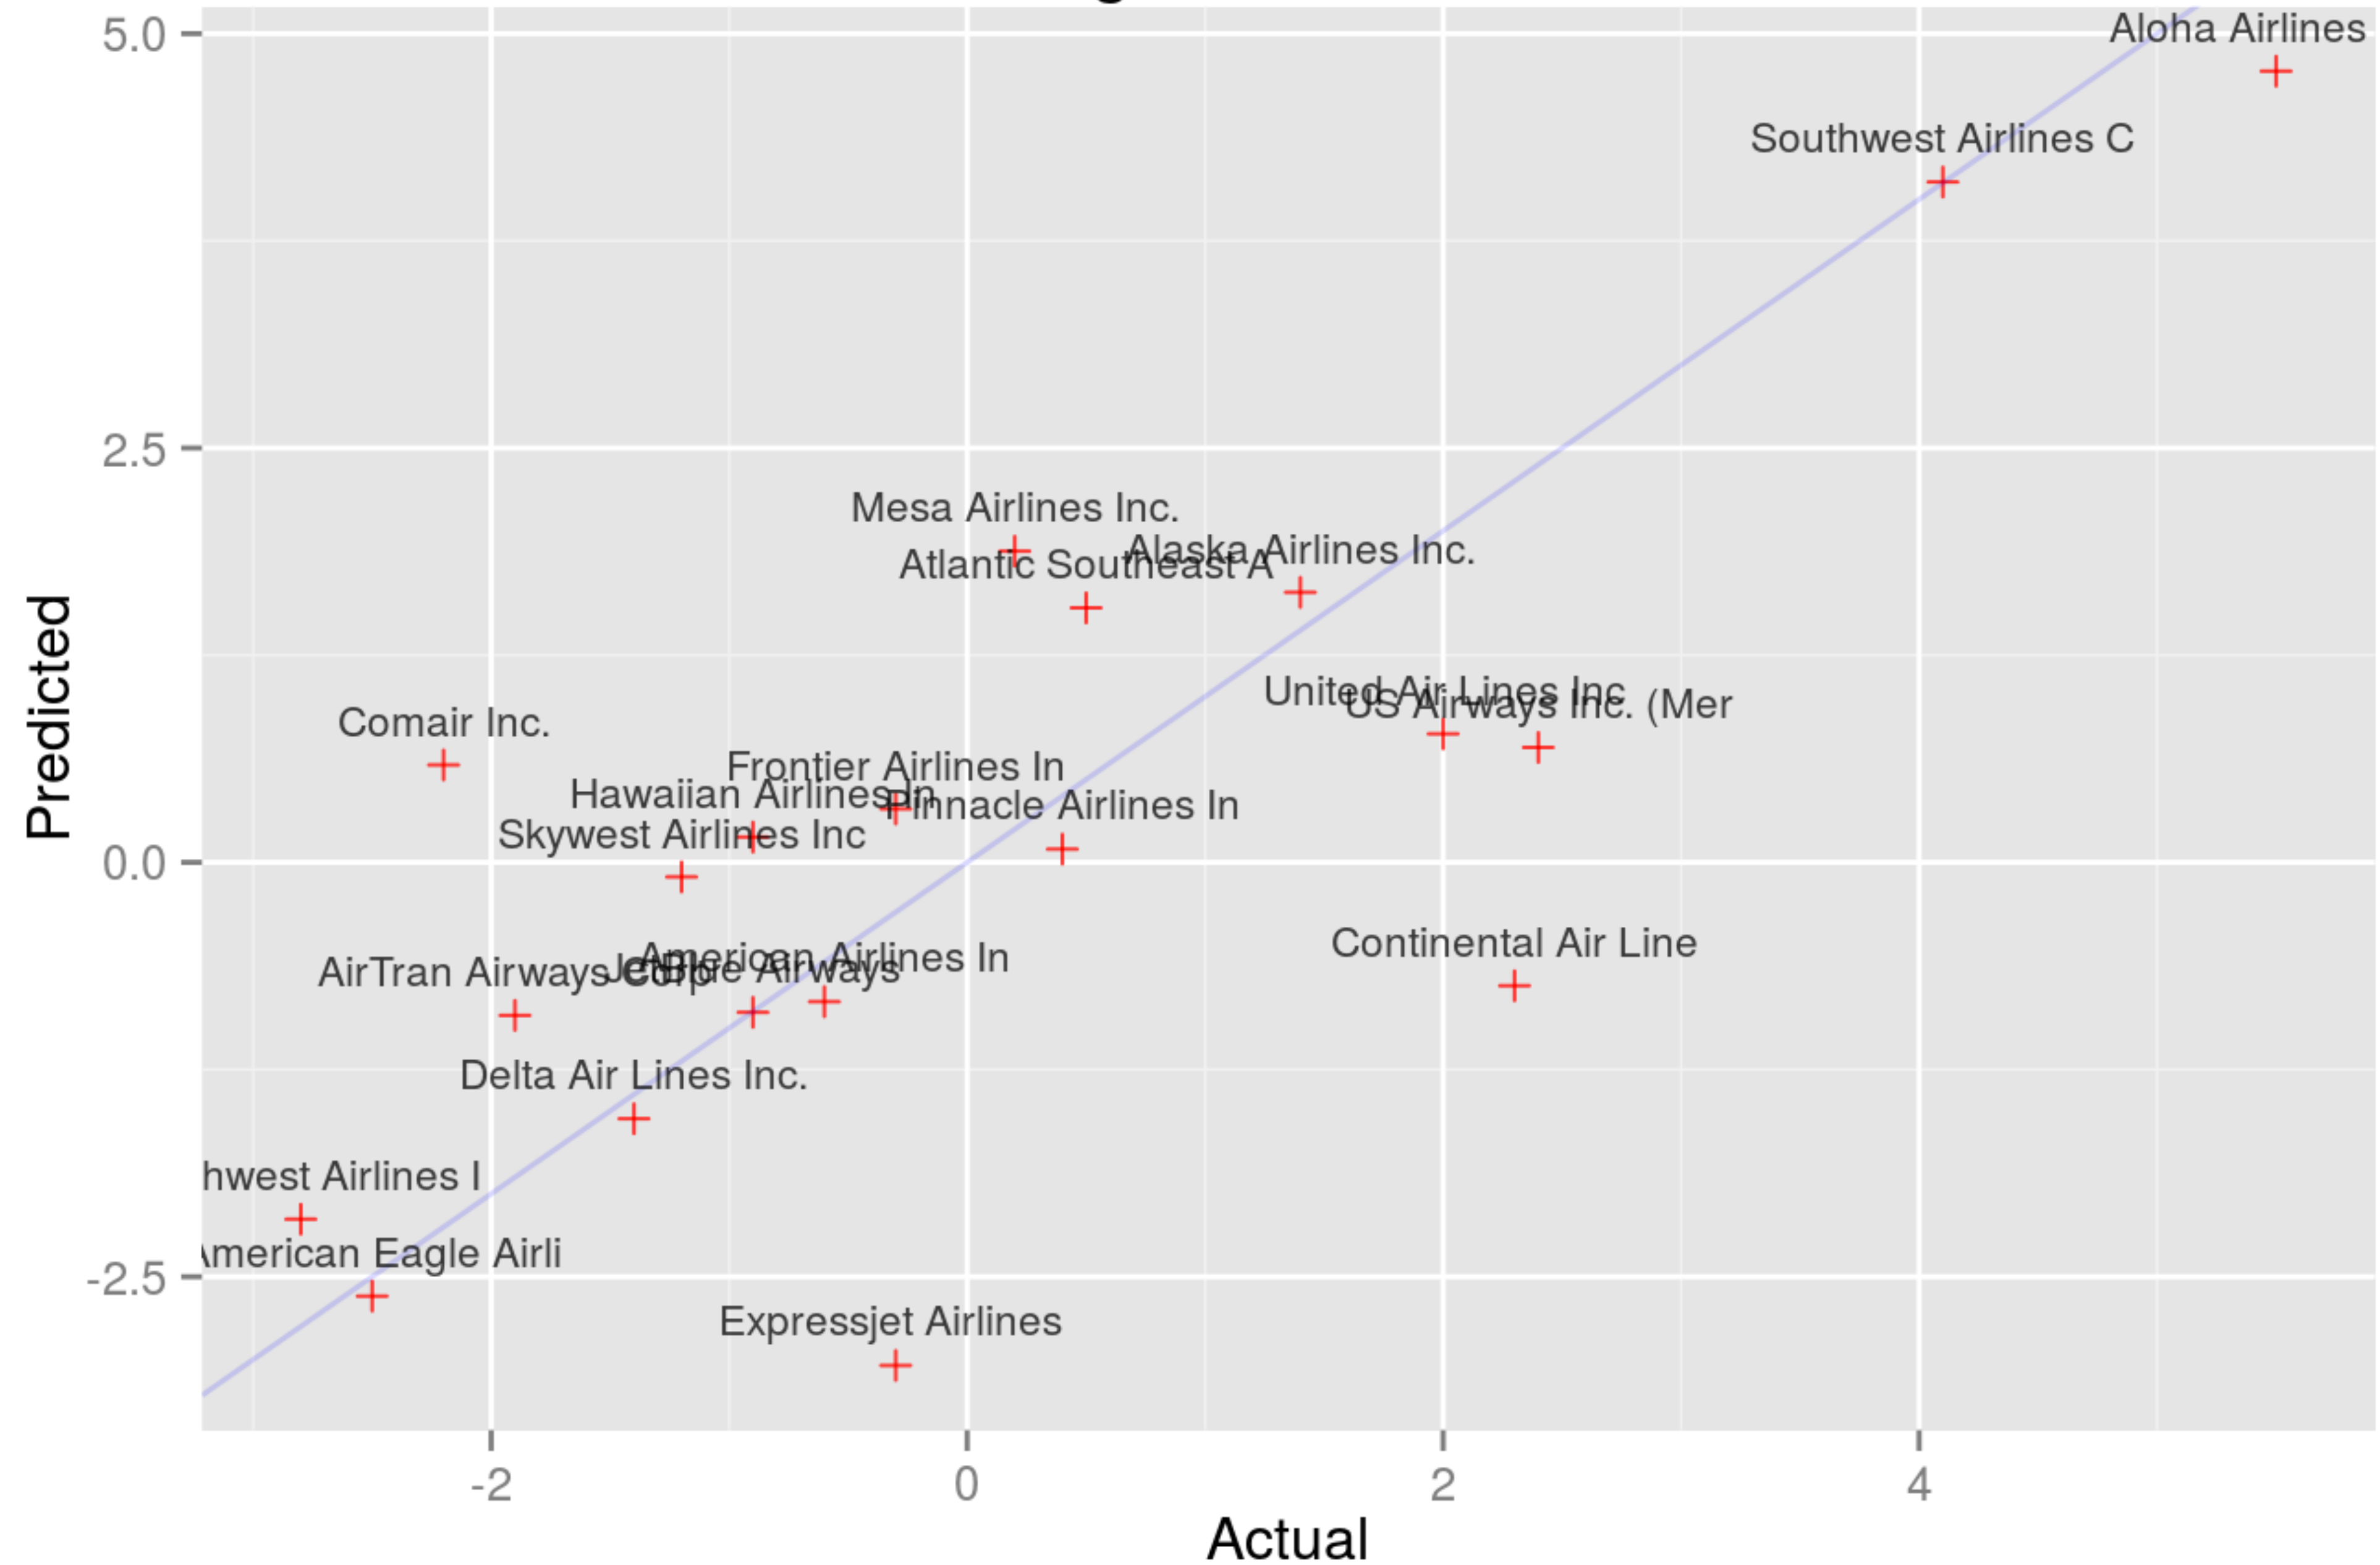
Score data (in DBMS)

```
# Score test data
score <- flights %>%
  filter(year == 2008) %>%
  filter(!is.na(arrdelay) & !is.na(depdelay) & !is.na(distance)) %>%
  filter(depdelay > 15 & depdelay < 240) %>%
  filter(arrdelay > -60 & arrdelay < 360) %>%
  select(arrdelay, depdelay, distance, uniquecarrier) %>%
  left_join(carriers, by = c('uniquecarrier' = 'code')) %>%
  left_join(coefs_table, copy = TRUE) %>%
  mutate(gain = depdelay - arrdelay) %>%
  mutate(pred = int_score + carrier_score + dist_score * distance +
           delay_score * depdelay) %>%
  group_by(description) %>%
  summarize(gain = mean(1.0 * gain), pred = mean(pred))
scores <- collect(score)
```

Visualize scores

```
library(ggplot2)
ggplot(scores, aes(gain, pred)) +
  geom_point(alpha = 0.75, color = 'red', shape = 3) +
  geom_abline(intercept = 0, slope = 1, alpha = 0.15, color = 'blue') +
  geom_text(aes(label = substr(description, 1, 20)),
    size = 4, alpha = 0.75, vjust = -1) +
  labs(title='Average Gains Forecast', x = 'Actual', y = 'Predicted')
```

Average Gains Forecast



copy_to

```
copy_to(air, query5, name = "gains")
```

copy_to()

connection
object

data frame to copy
to data database

table name in
database

`copy_to()` creates a table in the database from a local data frame.

Close Connection

```
rm(air)
```


Remove connection object

```
gc()
```

Run garbage collector

dplyr automatically closes connections when you remove the connection object *and then run the garbage collector, gc().*

Alternative APIs



RHadoop



SparkR



RevoScaleR



PivotalR



RODBC




RJDBC, etc...

Recap: Access Big Data



Store data outside of memory in data warehouse



Use an R package as an API to the data warehouse. dplyr, DBI, sparkR, others.

```
db <-
```

Create and work with connection object

```
rm(db)  
gc()
```


Close connection when finished

Applying **big data**

Big Data and Visualization

- 1.** Plot **summaries**, not raw data

Big data ♥'s overplotting



Big Data and Visualization


- 1.** Plot **summaries**, not raw data
- 2.** Calculate summaries **manually** in data store


```
cldata <- collect(clean)
ggplot(cldata) +
  geom_bar(aes(x = uniquecarrier))
```

year	arrdelay	depdelay	distance	unique carrier
2007	42	40	424	9E
2007	90	94	424	9E
2007	19	20	424	9E
2007	184	167	424	9E
2007	21	30	424	9E
2007	178	179	424	9E
2007	56	59	424	9E
2007	21	21	424	9E
2007	50	57	424	9E
...



unique carrier	count
DH	123752
US	431913
AA	677471
F9	37710
HP	105926
AS	189748
AQ	5368
9E	38367
EV	326694
NW	381213






```
cldata <- collect(clean)
ggplot(cldata) +
  geom_bar(aes(x = uniquecarrier),
    stat = "bin")
```

year	arrdelay	depdelay	distance	unique carrier
2007	42	40	424	9E
2007	90	94	424	9E
2007	19	20	424	9E
2007	184	167	424	9E
2007	21	30	424	9E
2007	178	179	424	9E
2007	56	59	424	9E
2007	21	21	424	9E
2007	50	57	424	9E
...




unique carrier	count
DH	123752
US	431913
AA	677471
F9	37710
HP	105926
AS	189748
AQ	5368
9E	38367
EV	326694
NW	381213



```
clsummary <- clean %>%
  group_by(uniquecarrier) %>%
  summarise(count = n()) %>%
  collect()
```

unique carrier	count
DH	123752
US	431913
AA	677471
F9	37710
HP	105926
AS	189748
AQ	5368
9E	38367
EV	326694
NW	381213

```
ggplot(clsummary) +
  geom_bar(aes(x = uniquecarrier,
               y = count),
           stat = "identity")
```



R Markdown **and big data**

Big Data and R Markdown

- 1.** **Cache** code chunks that manipulate big data.

cache

R Markdown will cache the result of the code chunk to reuse (and thus avoid computation) when **cache = TRUE**

Here's some code that takes a "long" time to run.

```
```{r cache=TRUE}
```

```
Sys.sleep(5)
```

```
rnorm(1)
```

```
```
```

Here's some code that takes a "long" time to run.

```
Sys.sleep(5)  
rnorm(1)
```

```
## [1] 1.582407
```

Here's some code that takes a "long" time to run.

```
Sys.sleep(5)  
rnorm(1)
```

```
## [1] 1.582407
```

cache

R Markdown will cache the result of the code chunk to reuse (and thus avoid computation) when **cache = TRUE**

```
```{r cache=TRUE}  
d <- flights %>%
 select(...) %>%
 filter(...) %>%
 mutate(...) %>%
 collect()
```
```

Here's some code that
takes a "long" time to run.

```
Sys.sleep(5)  
rnorm(1)
```

```
## [1] 1.582407
```

Here's some code that
takes a "long" time to run.

```
Sys.sleep(5)  
rnorm(1)
```

```
## [1] 1.582407
```



Big Data and R Markdown

- 1.** **Cache** code chunks that manipulate big data.
- 2.** Set **engine** to use another language if sensible.

engine

To embed non R code, set the engine option to the language you want to embed.

```
Some python code,  
```{r engine='python'}  
x = 'hello, python
world!'
print(x)
print(x.split(' '))
```
```



Some python code:

```
x = 'hello, python world!'  
print(x)  
print(x.split(' '))
```

```
## hello, python world!  
## ['hello,', 'python', 'world!']
```

knitr comes with engines for the following languages, and can be extended to other languages

asis

asy

awk

bash

c

cat

coffee

dot

fortran

gawk

groovy

haskell

highlight

lein

mysql

node

perl

psql

python

Rcpp

Rscript

ruby

sas

scala

sed

sh

stan

stata


tikz


zsh

Shiny **and big data**

Big Data and Shiny

- 1.** **Avoid** unnecessary repetitions



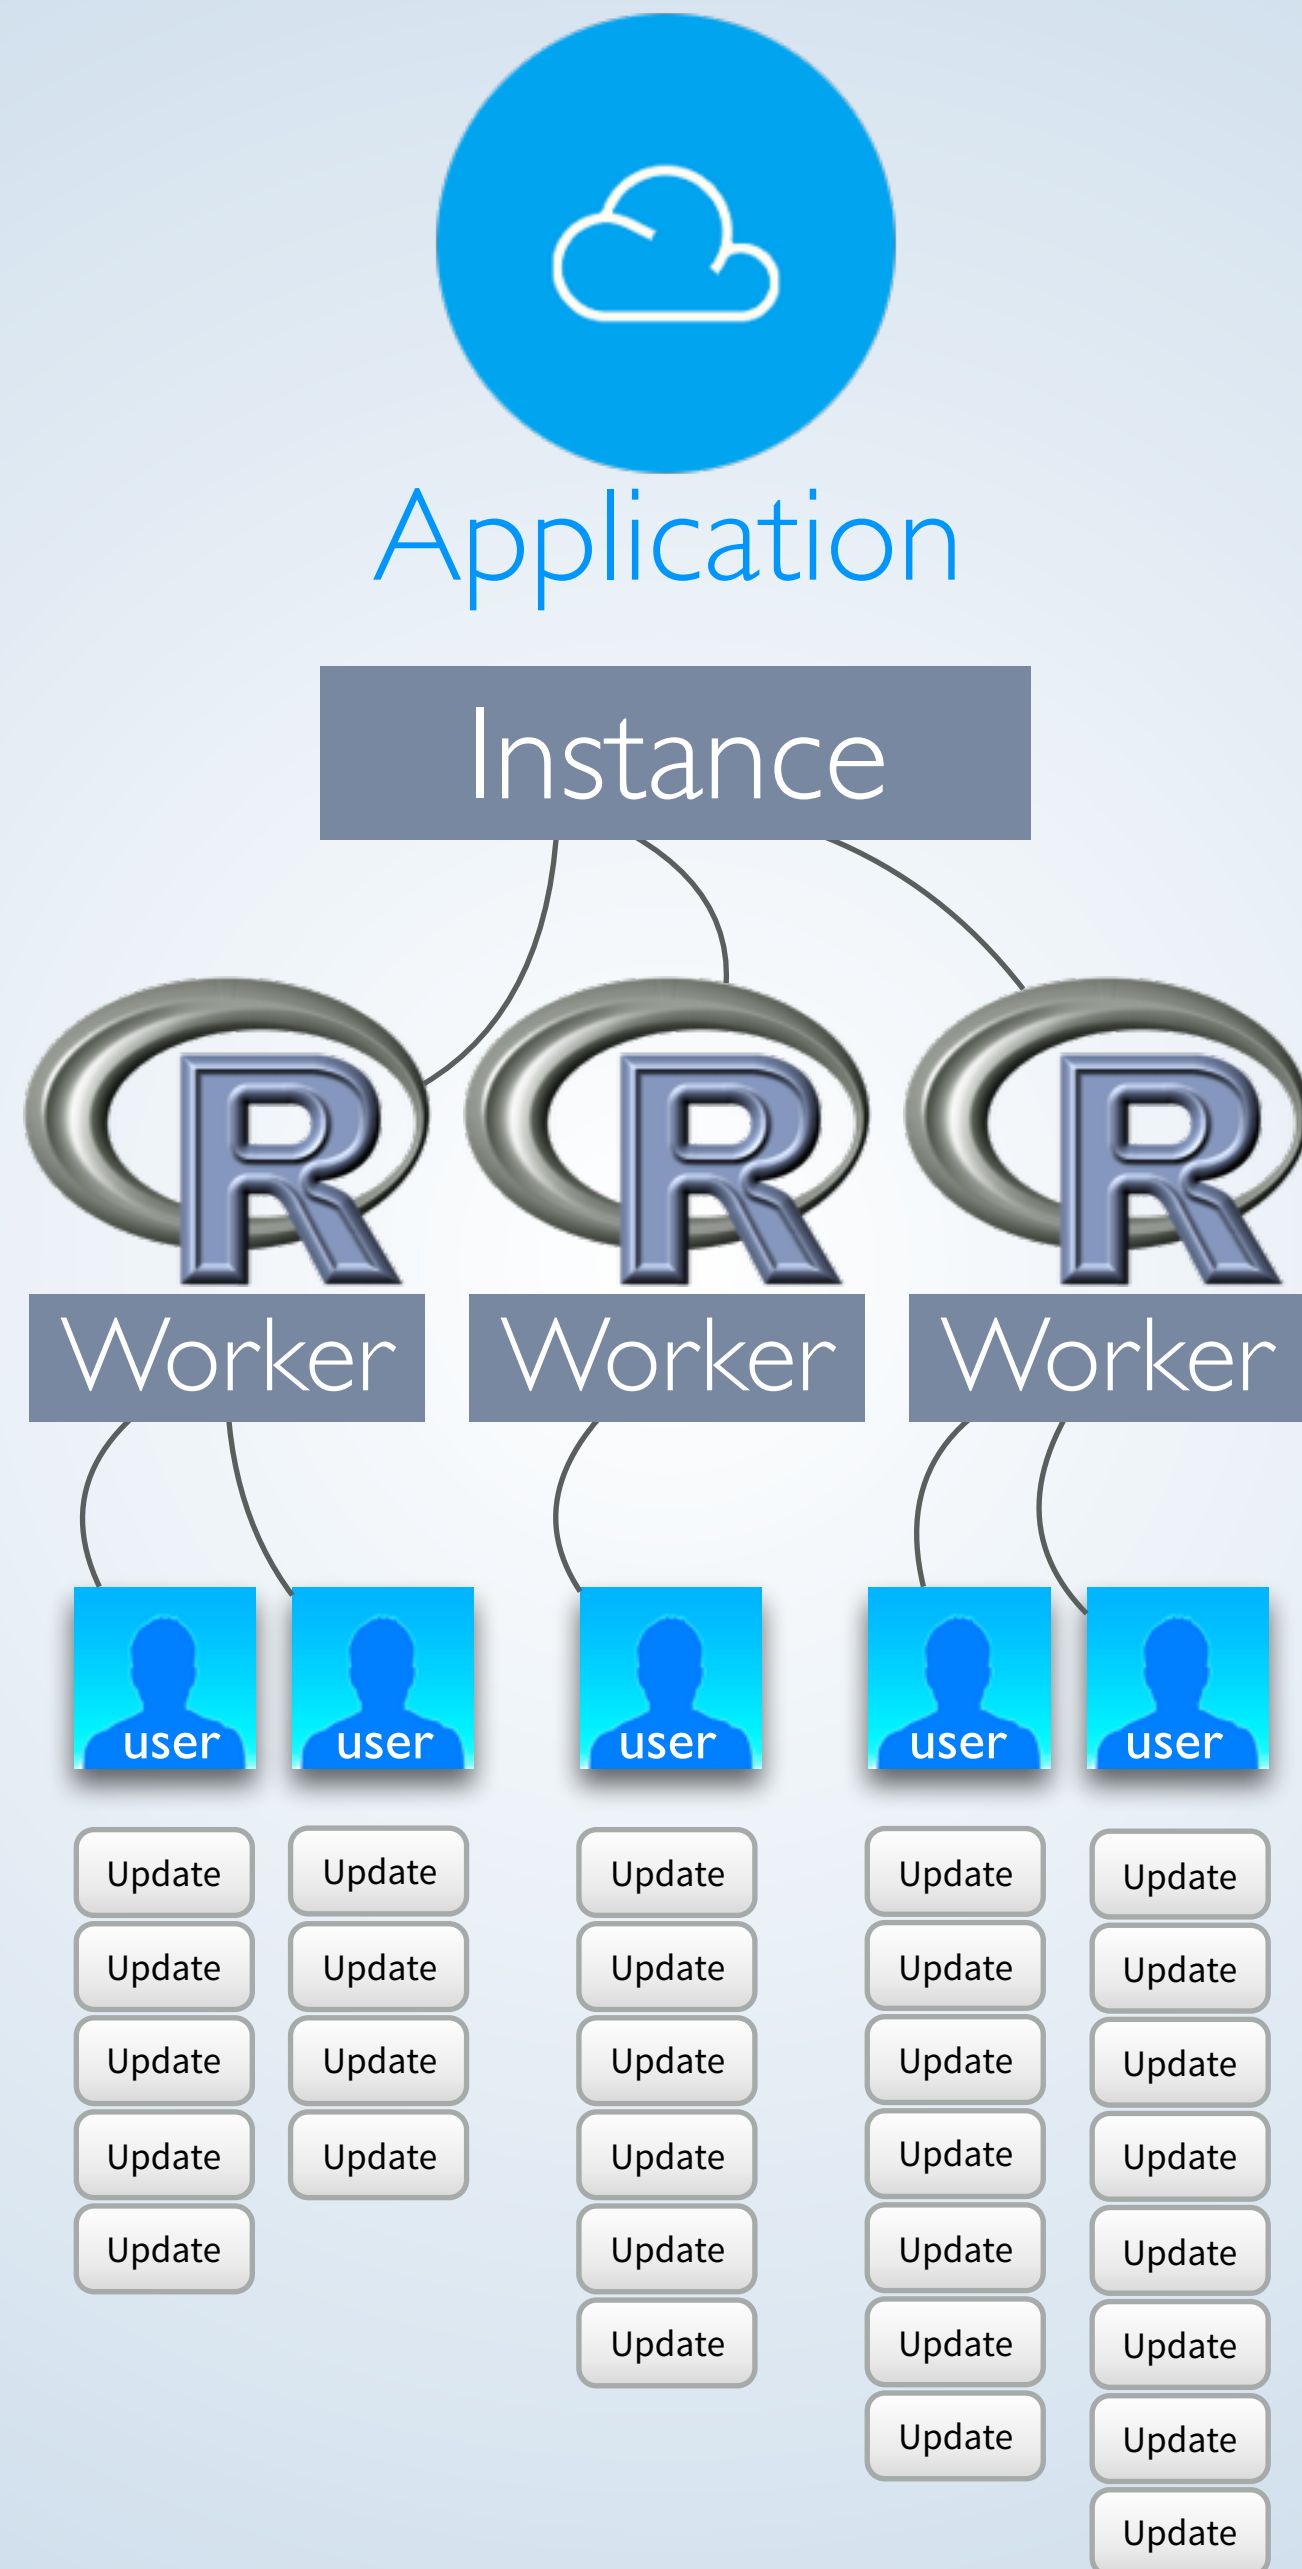


```
library(shiny)
```

```
ui <- fluidPage(  
  sliderInput(inputId = "num",  
    label = "Choose a number",  
    value = 25, min = 1,  
    max = 100),  
  plotOutput("hist")  
)
```

```
server <- function(input, output) {  
  output$hist <- renderPlot({  
    hist(rnorm(input$num))  
  })  
}
```

```
shinyApp(ui = ui, server = server)
```




Code outside the server function will be run once per R worker

```
library(shiny)
```

```
ui <- fluidPage(  
  sliderInput(inputId = "num",  
    label = "Choose a number",  
    value = 25, min = 1,  
    max = 100),  
  plotOutput("hist")  
)
```

```
server <- function(input, output) {  
  output$hist <- renderPlot({  
    hist(rnorm(input$num))  
  })  
}
```

```
shinyApp(ui = ui, server = server)
```



Code outside the server function will be run once per R worker

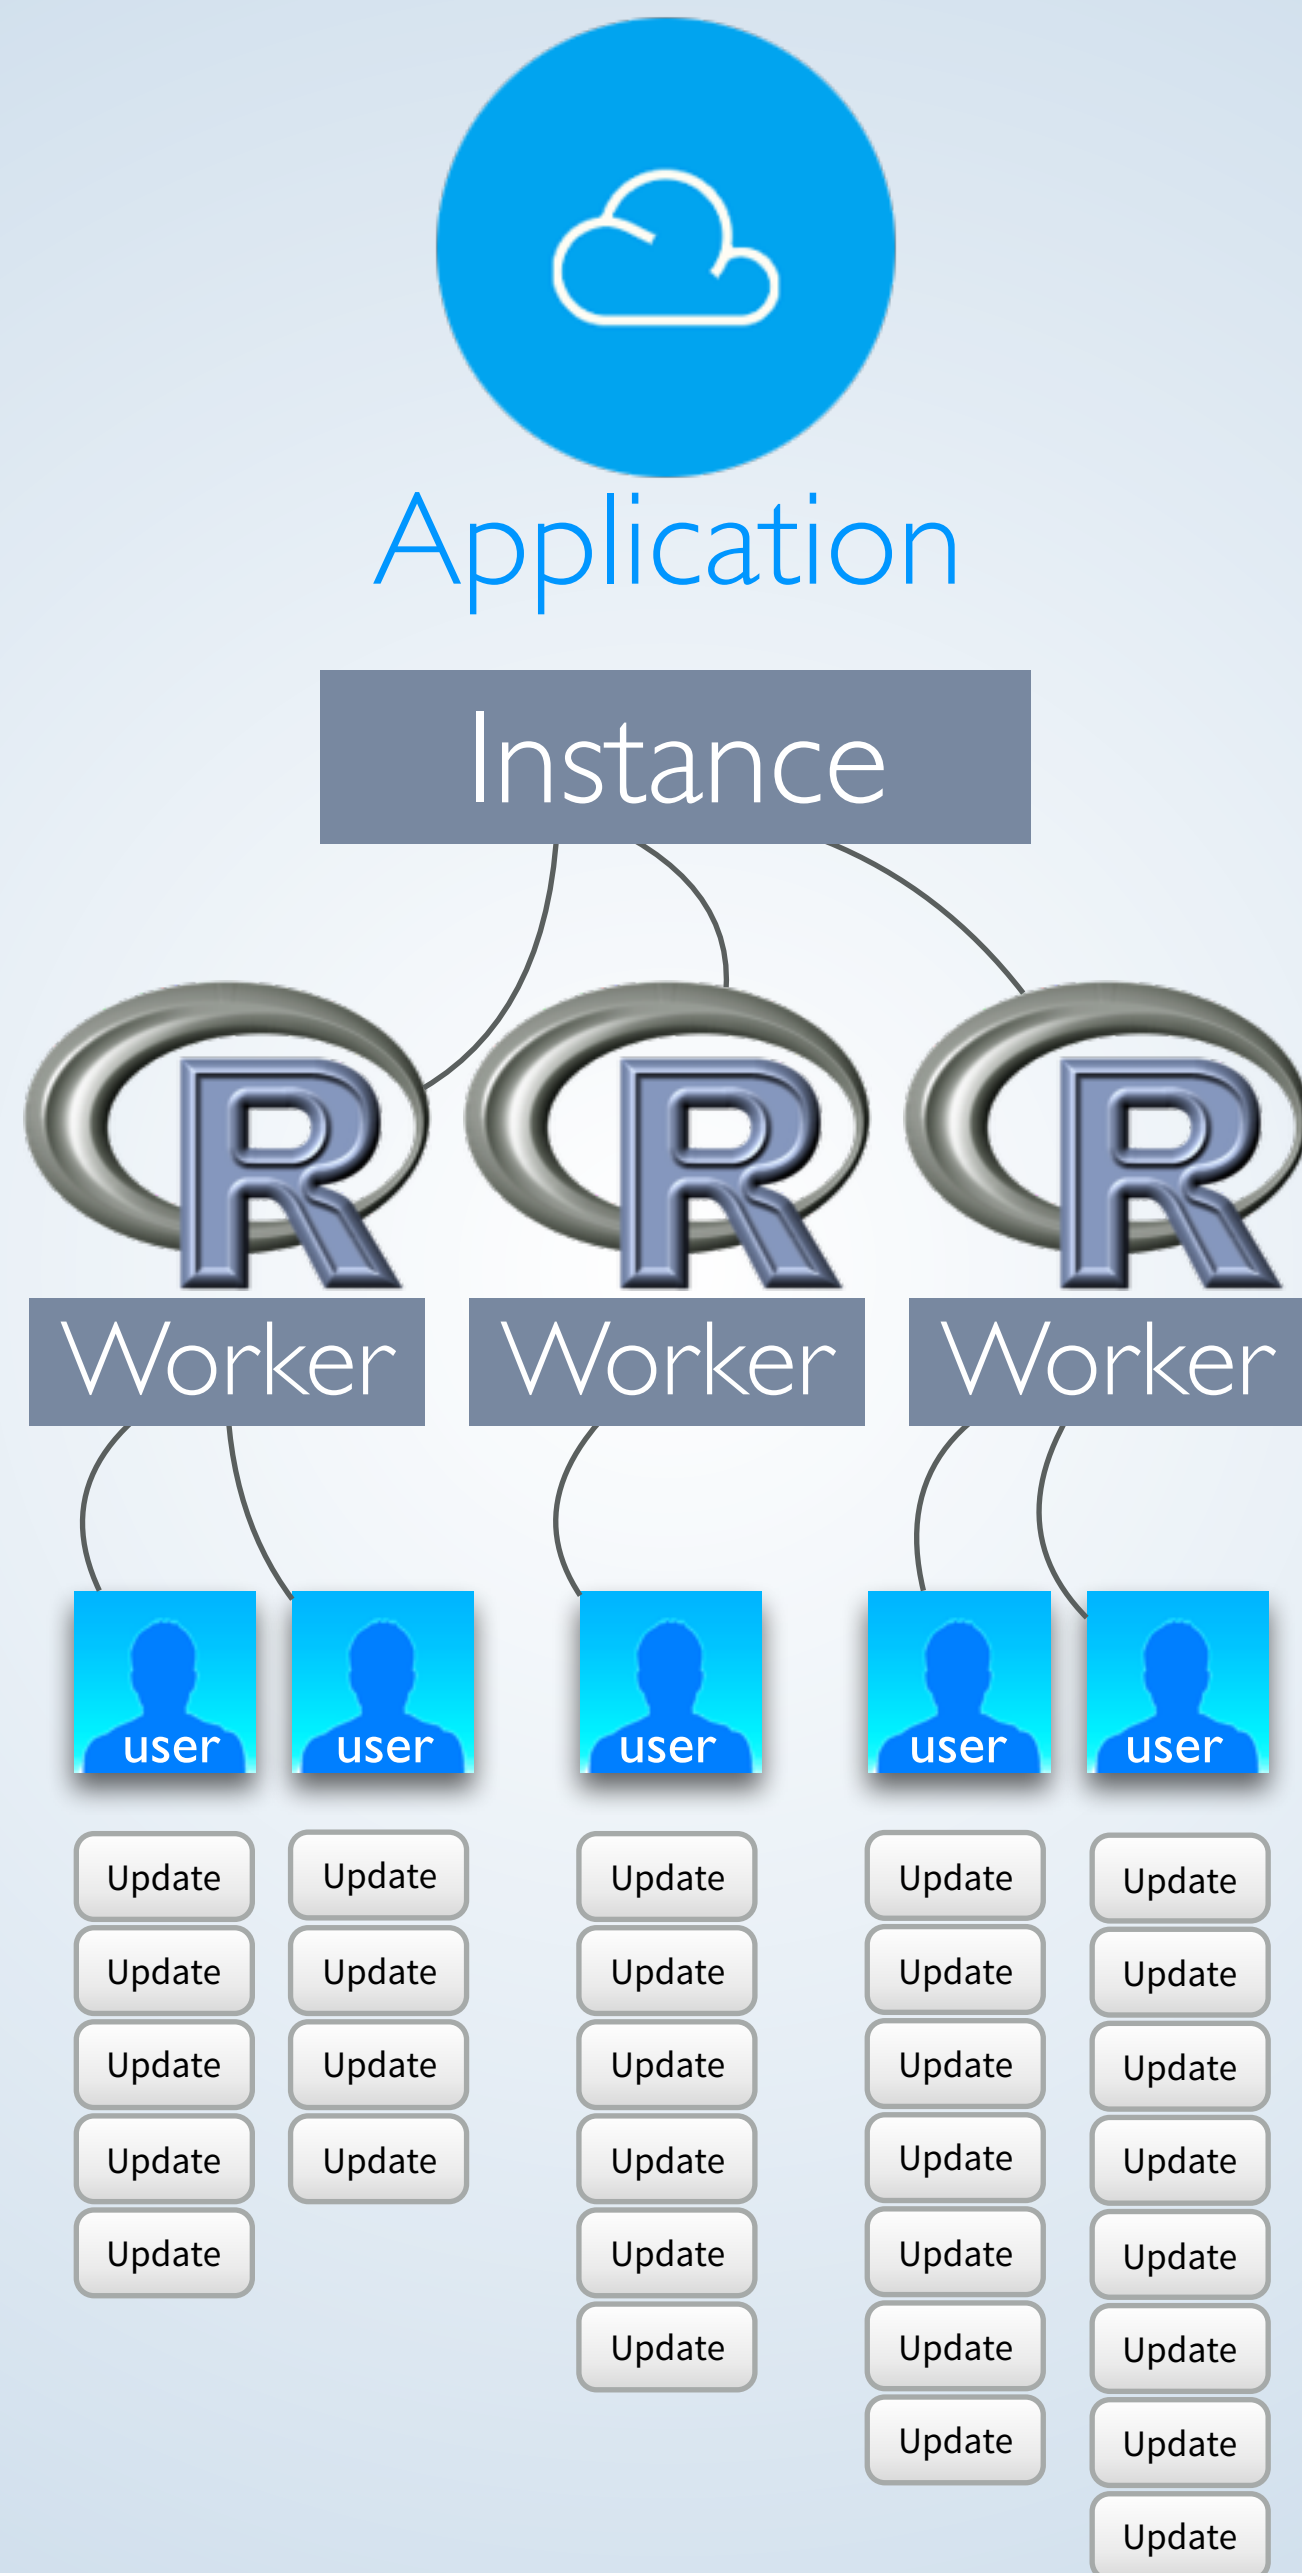
Code inside the server function will be run once per connection

```
library(shiny)
```

```
ui <- fluidPage(  
  sliderInput(inputId = "num",  
    label = "Choose a number",  
    value = 25, min = 1,  
    max = 100),  
  plotOutput("hist")  
)
```

```
server <- function(input, output) {  
  output$hist <- renderPlot({  
    hist(rnorm(input$num))  
  })  
}
```

```
shinyApp(ui = ui, server = server)
```



Code outside the server function will be run once per R worker

Code inside the server function will be run once per connection

Code inside of a reactive function will be run once per reaction

```
library(shiny)
```

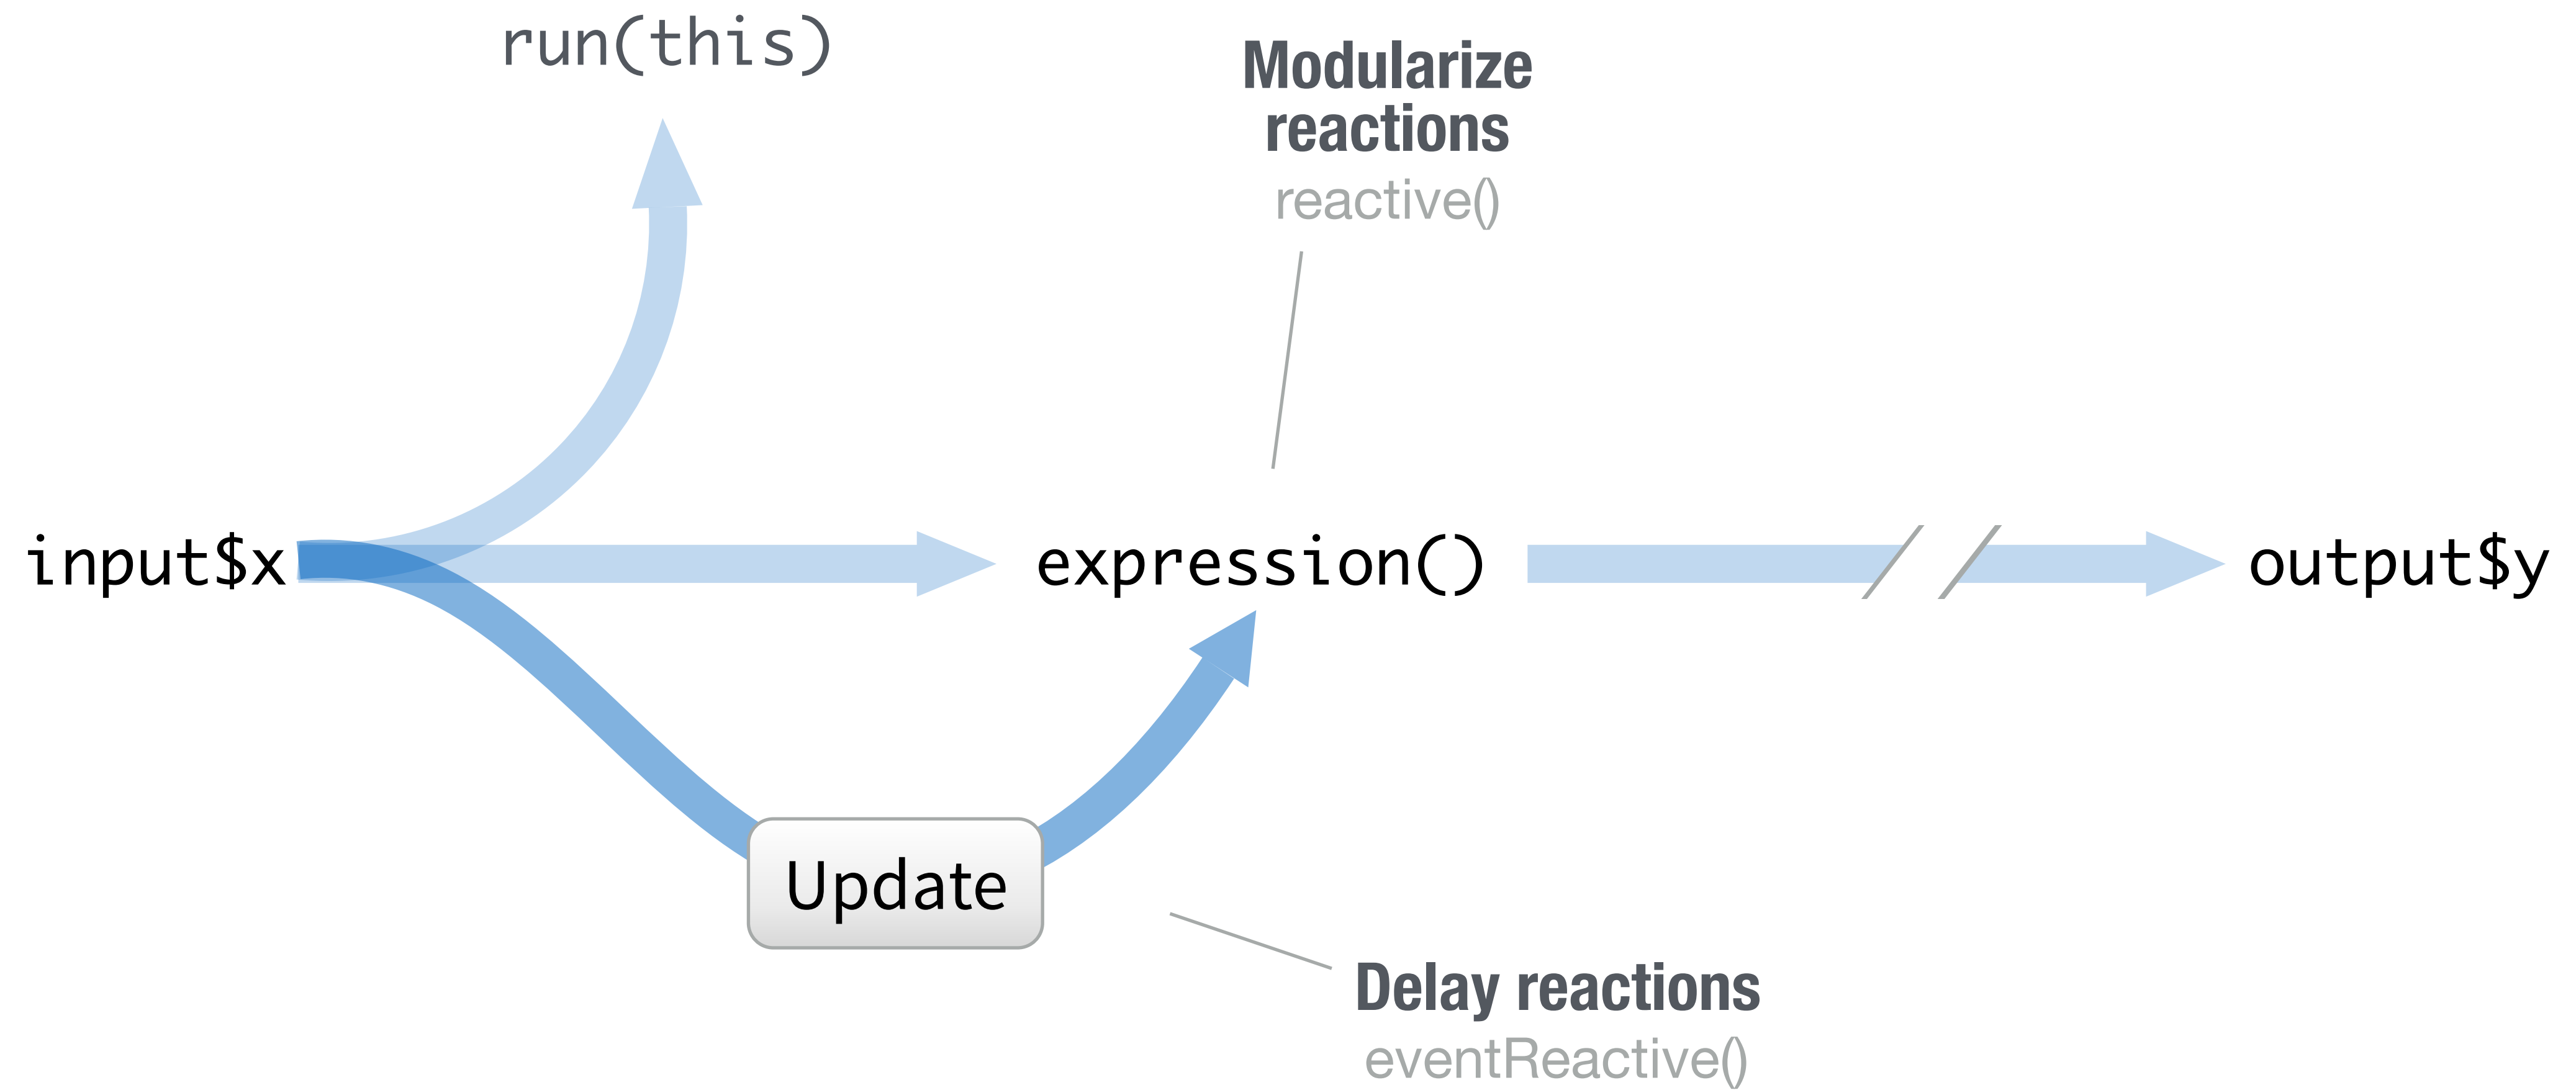
```
ui <- fluidPage(  
  sliderInput(inputId = "num",  
    label = "Choose a number",  
    value = 25, min = 1,  
    max = 100),  
  plotOutput("hist")  
)
```

```
server <- function(input, output) {  
  output$hist <- renderPlot({  
    hist(rnorm(input$num))  
  })  
}
```

```
shinyApp(ui = ui, server = server)
```


Big Data and Shiny

1. **Avoid** unnecessary repetitions
2. **Cache** expensive operations with reactive expressions
3. **Delay** expensive operations, e.g. with an action button



Thank You

